

# Cloud Computing

**Caltech**

**Center for Technology &  
Management Education**

## **Designing Infrastructure Solutions on Azure**



## Design a Solution for Relational Data

# A Day in the Life of an Azure Architect

You are an Azure architect for a company that needs a fully managed platform as a service (PaaS) database engine that can handle most database maintenance tasks like updating, patching, backups, and monitoring without requiring user interaction. They also want to make sure that the data is secure. They are also on the lookout for a relational database engine that is optimized for IoT and IoT Edge installations. The company is searching for a solution that can stream, process, and analyze relational and non-relational data.

Finally, your organization has asked you to help them with a fully managed NoSQL database service for modern app development.

To achieve all of the above, along with some additional features, we will be learning a few concepts in this lesson that will help you find a solution to the given scenario.



# Learning Objectives

By the end of this lesson, you will be able to:

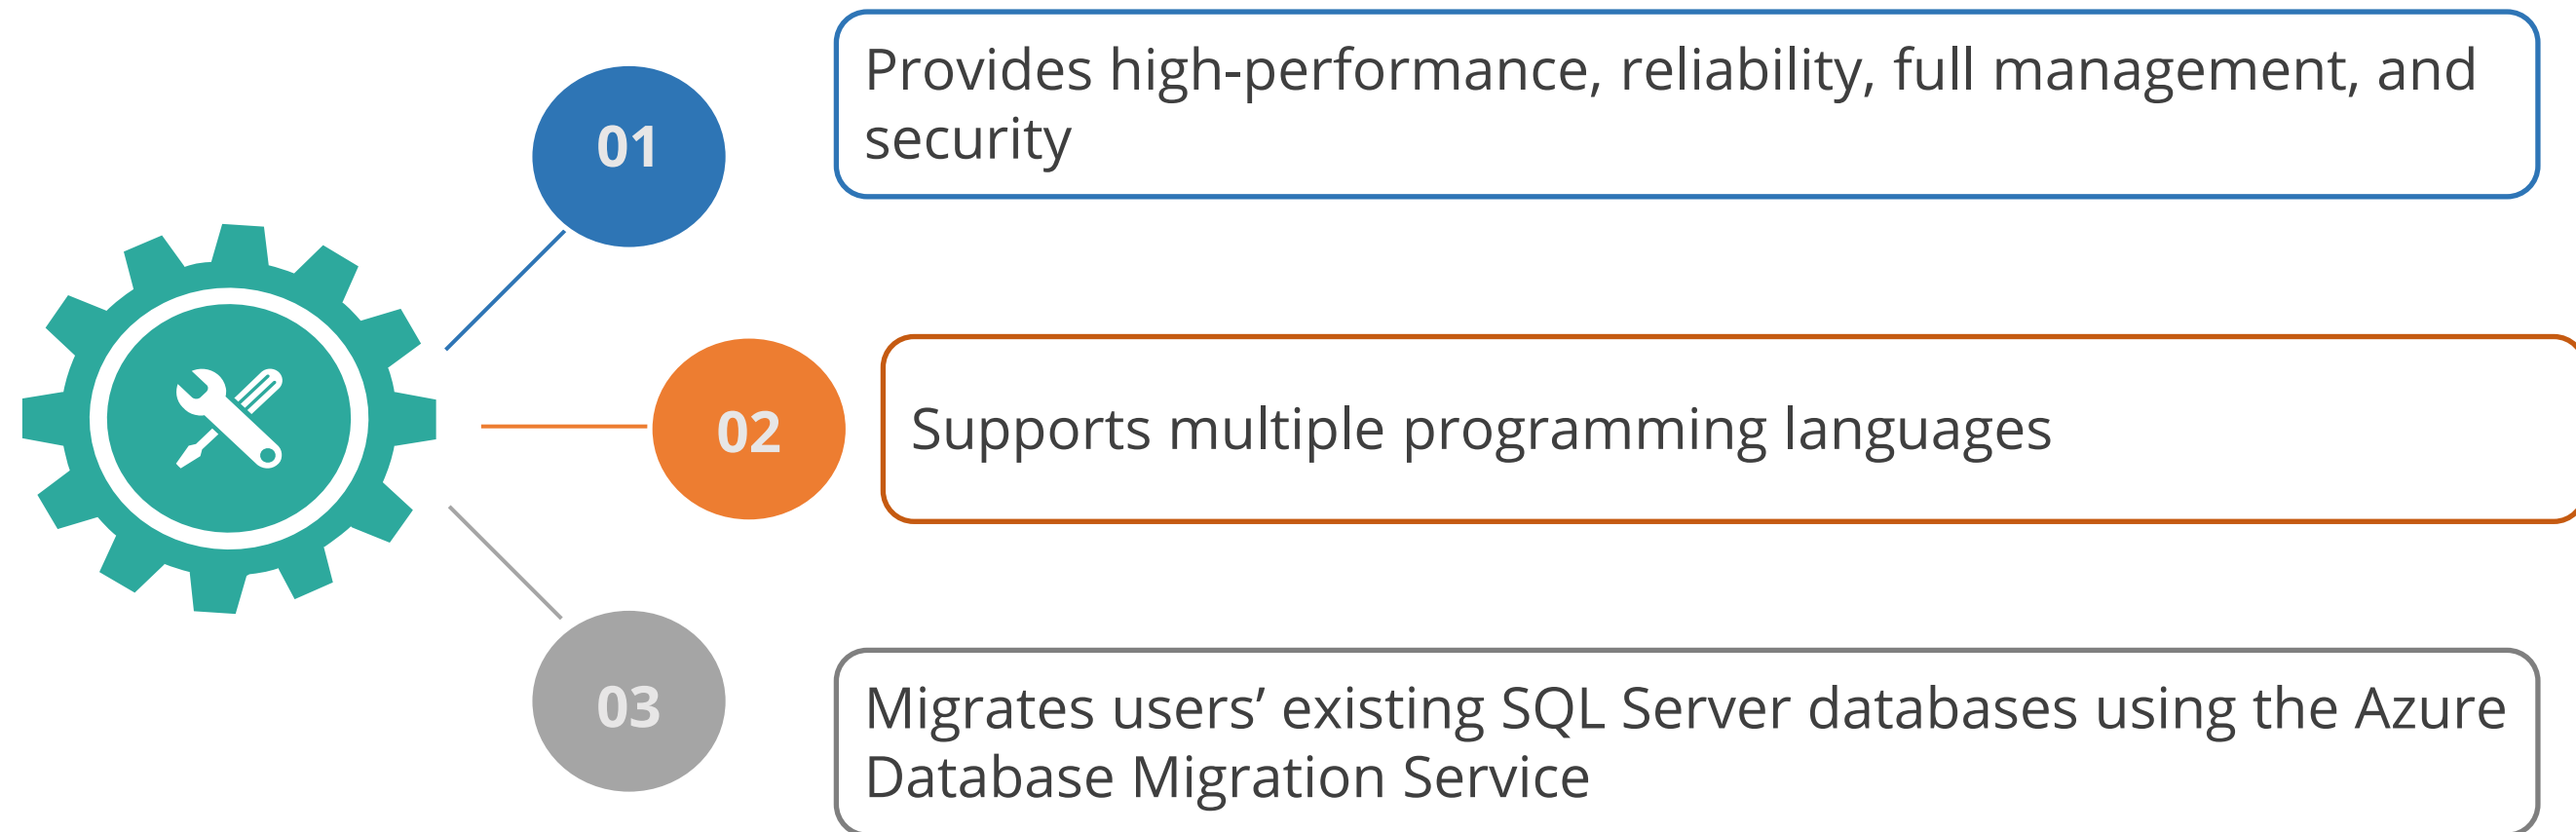
- Describe Azure SQL database
- Analyze an appropriate data platform based on requirements
- Illustrate Azure data storage
- Design for Azure SQL Edge and Azure Cosmos DB
- Implement design security for data



# Select an Appropriate Data Platform Based on Requirements

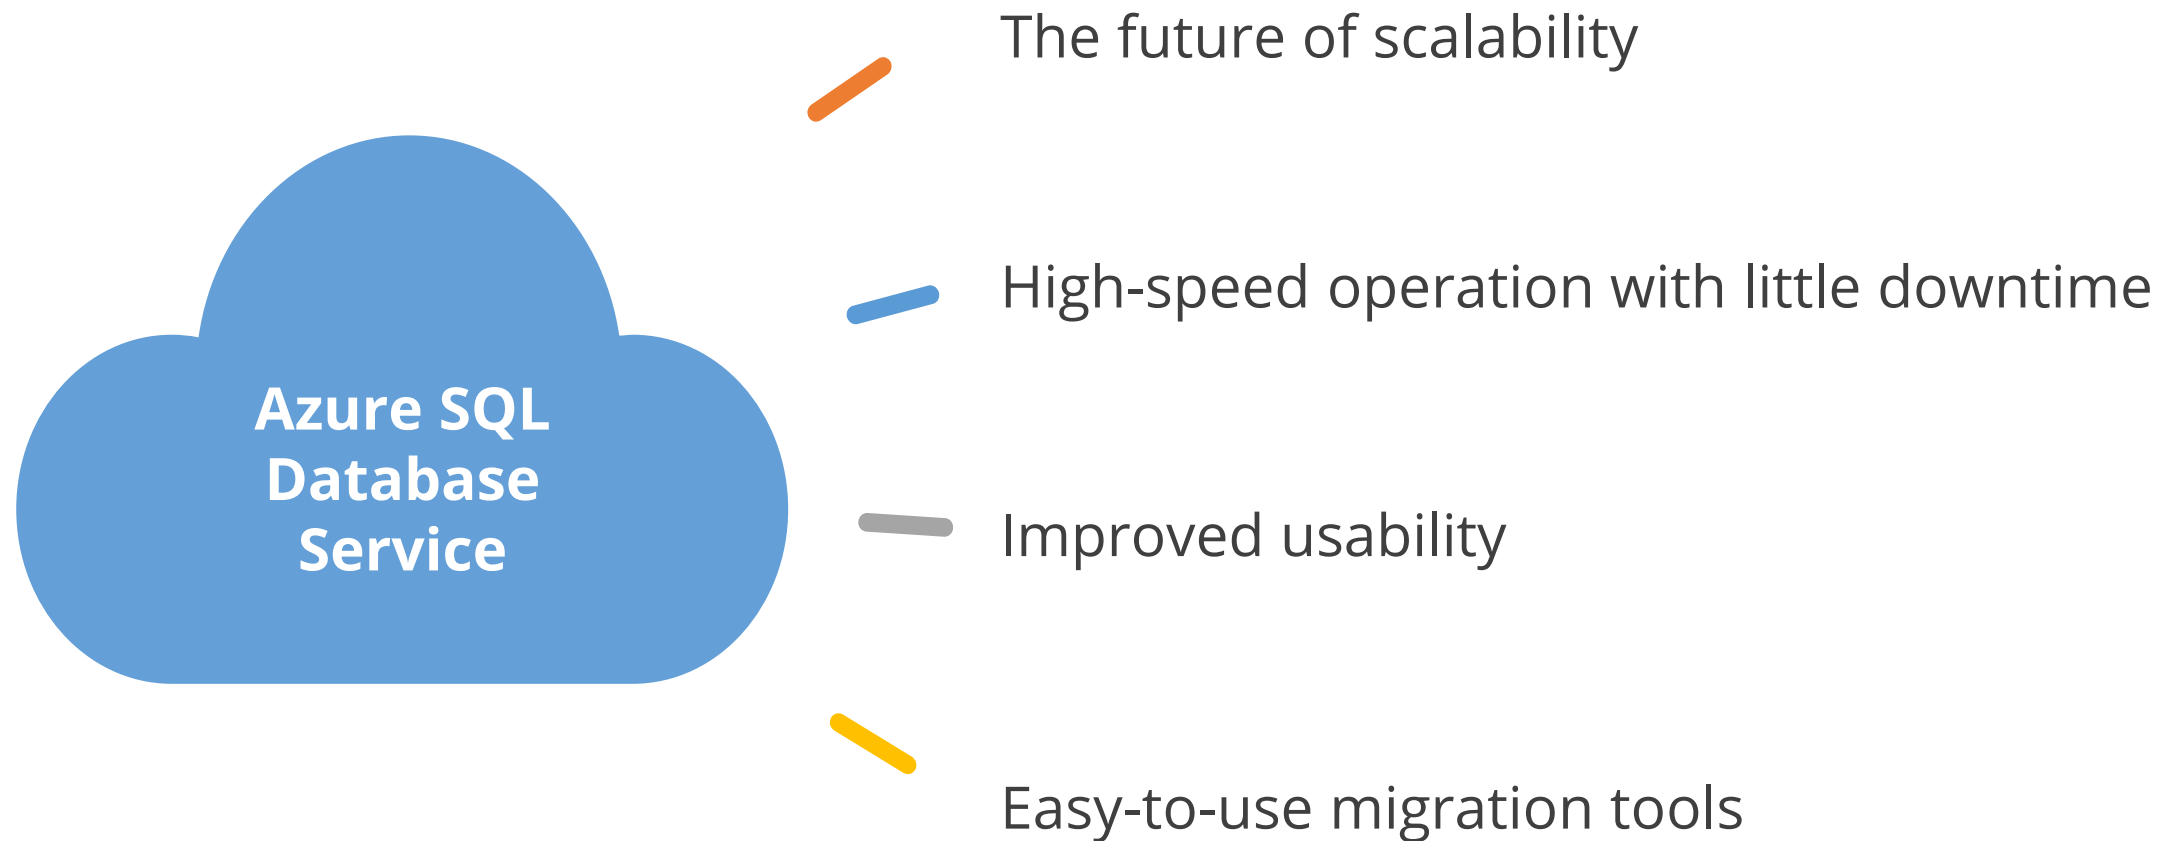
# Azure SQL Database

Azure SQL Database is a fully managed platform as a service (PaaS) database engine.



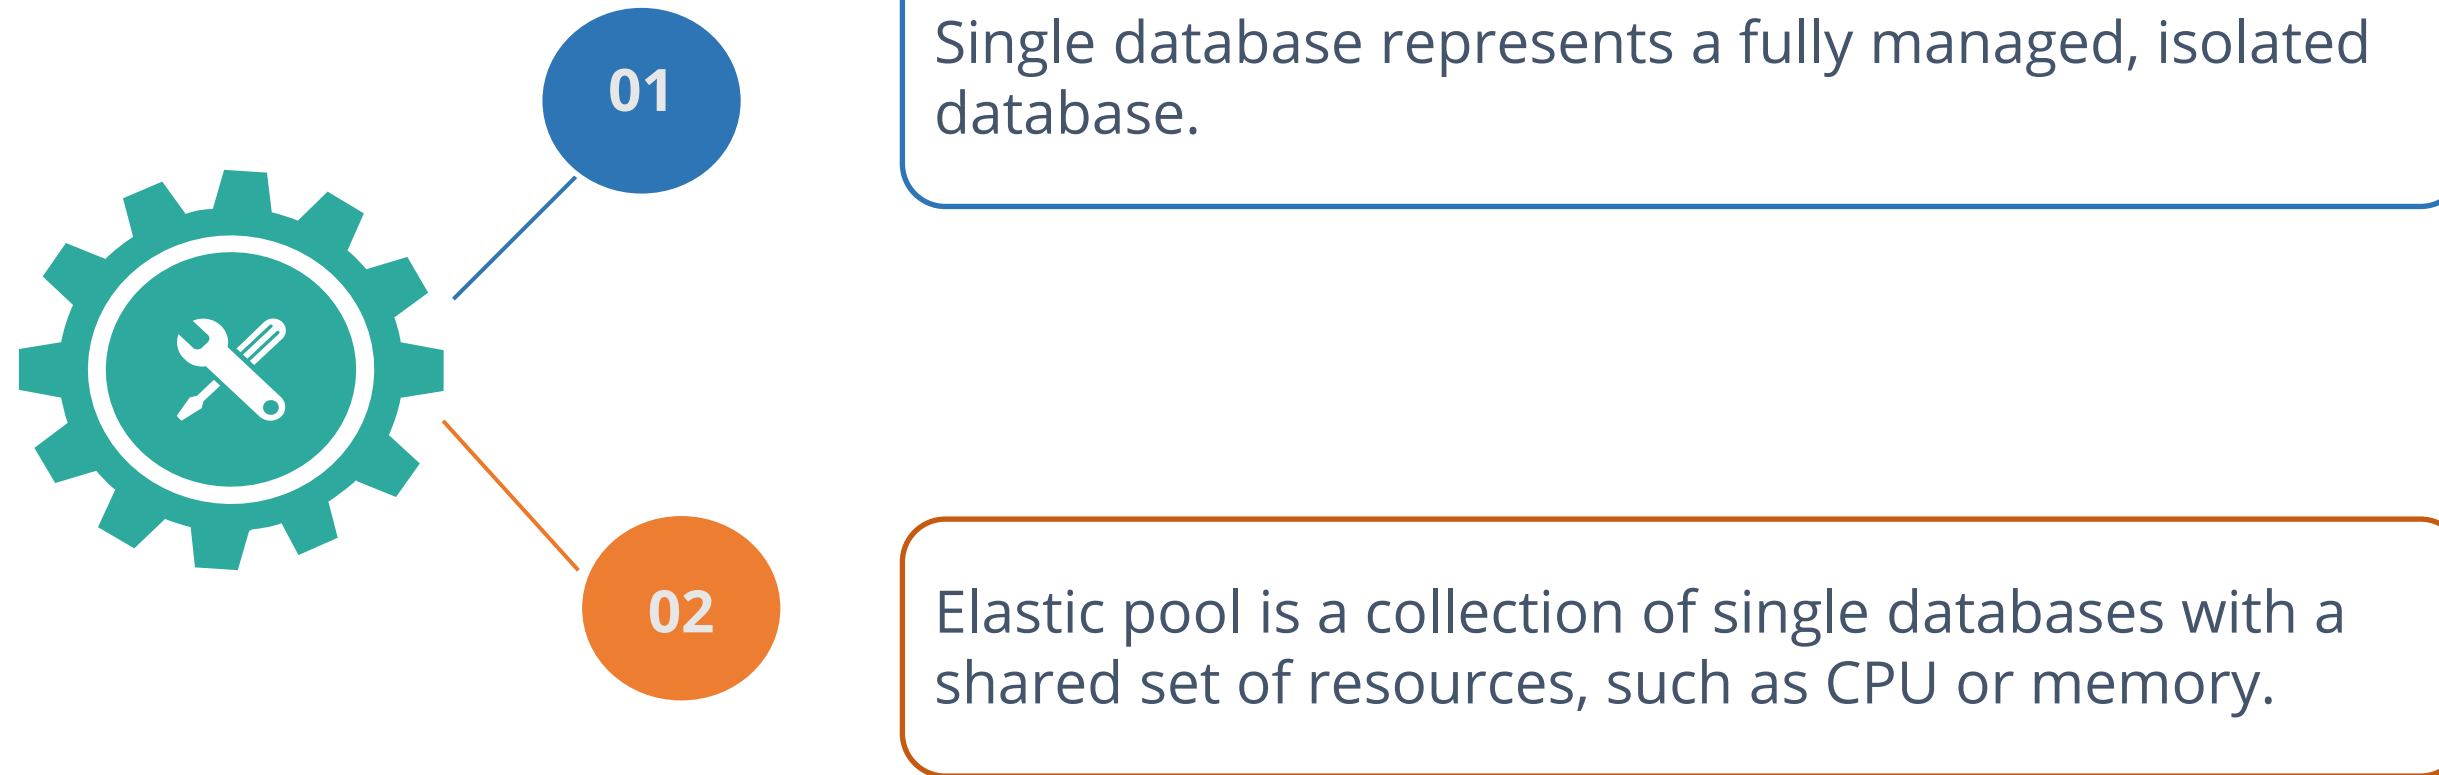
# Azure SQL Database Service

The benefits of Azure SQL Database Service are:



# Azure SQL Database

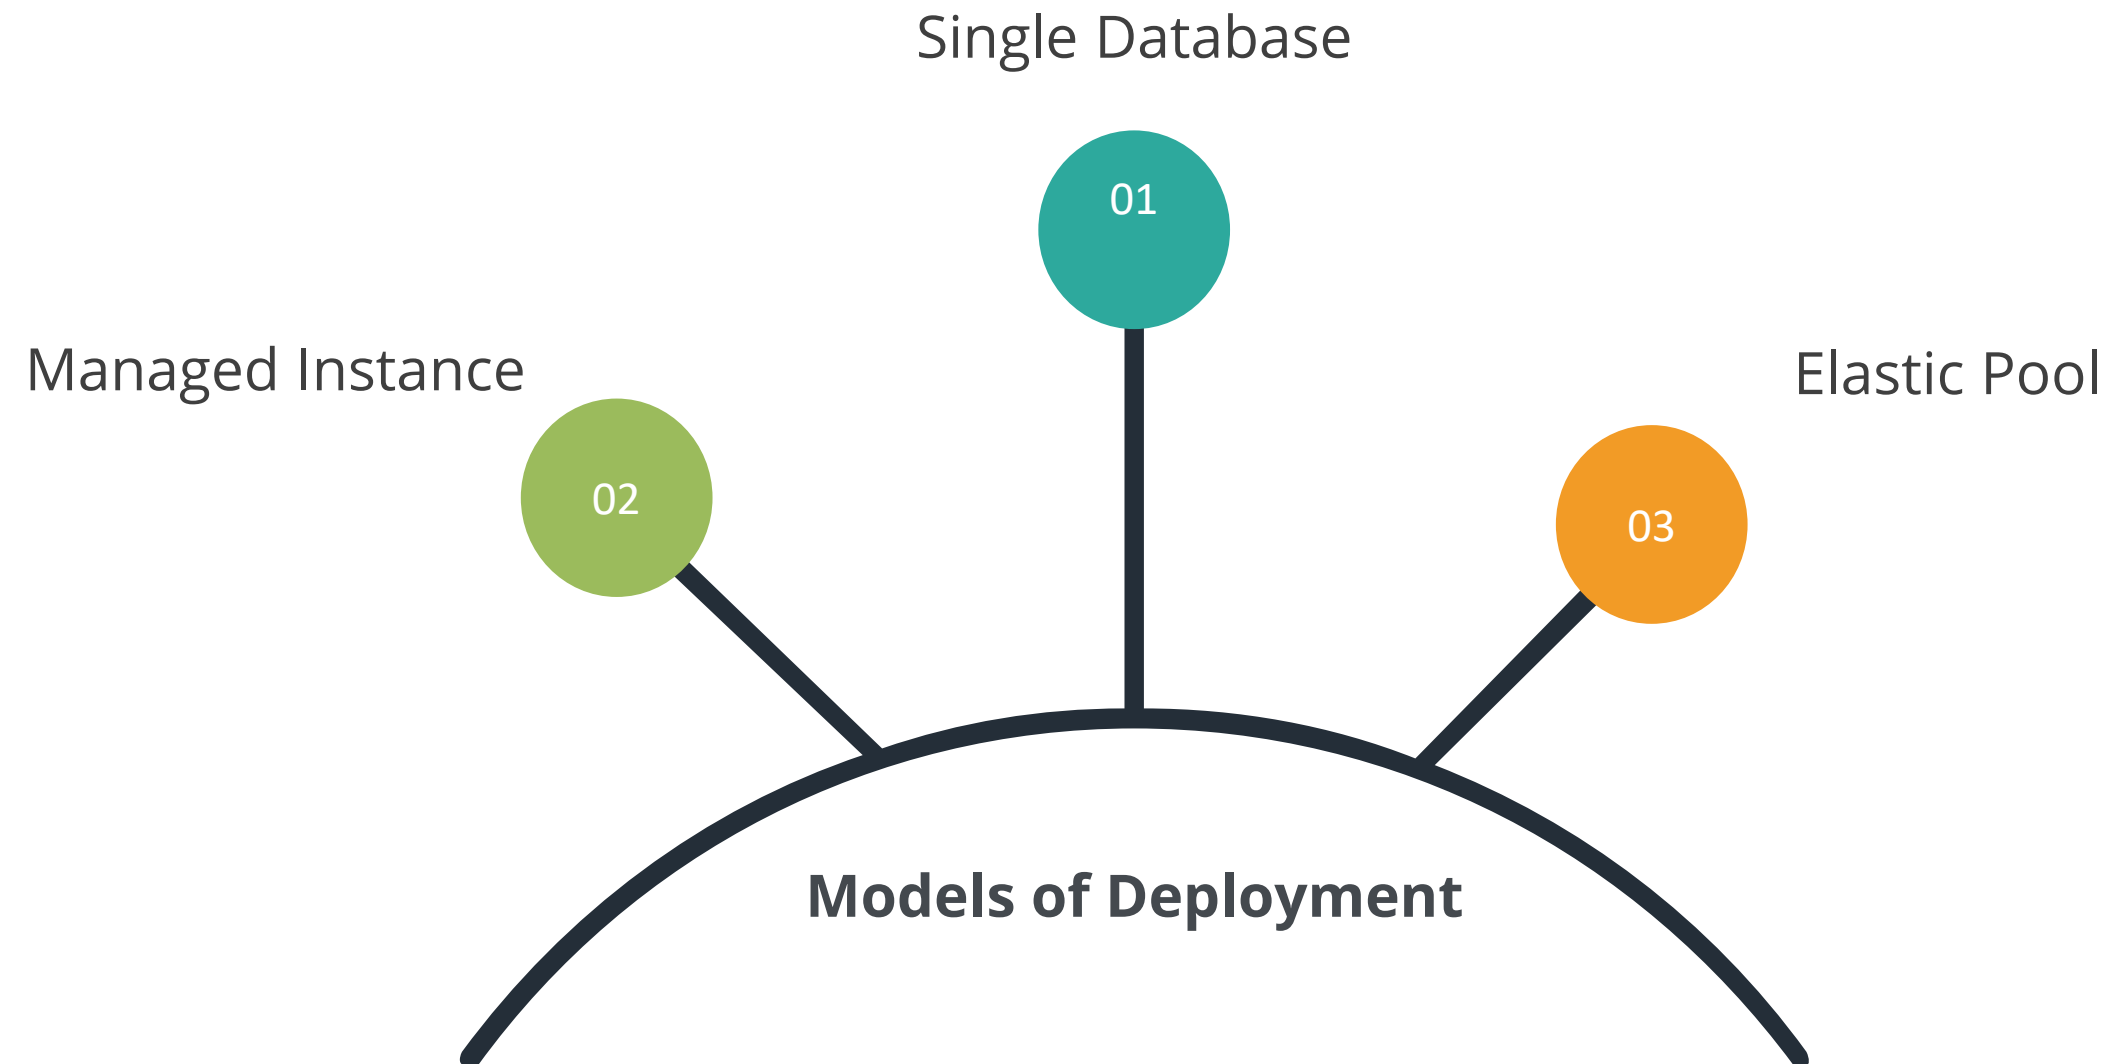
Azure SQL Database deployment options for a database:





# Deployment Models

There are three types of deployment models:



# Deployment Models

## Single Database

A single database is maintained by an SQL Database server and deployed on an Azure VM, and this is the simplest deployment method.

## Managed Instance

It is a database instance that is fully managed, which makes migrating on-premise SQL databases simple.

## Elastic Pool

It is a collection of interconnected databases that pool resources, and individual databases can be added and removed from an elastic pool.

# Deploy Azure SQL Database

There are five sections in the Azure portal to filled out during a deployment for Azure SQL Database and Azure SQL Managed Instance:

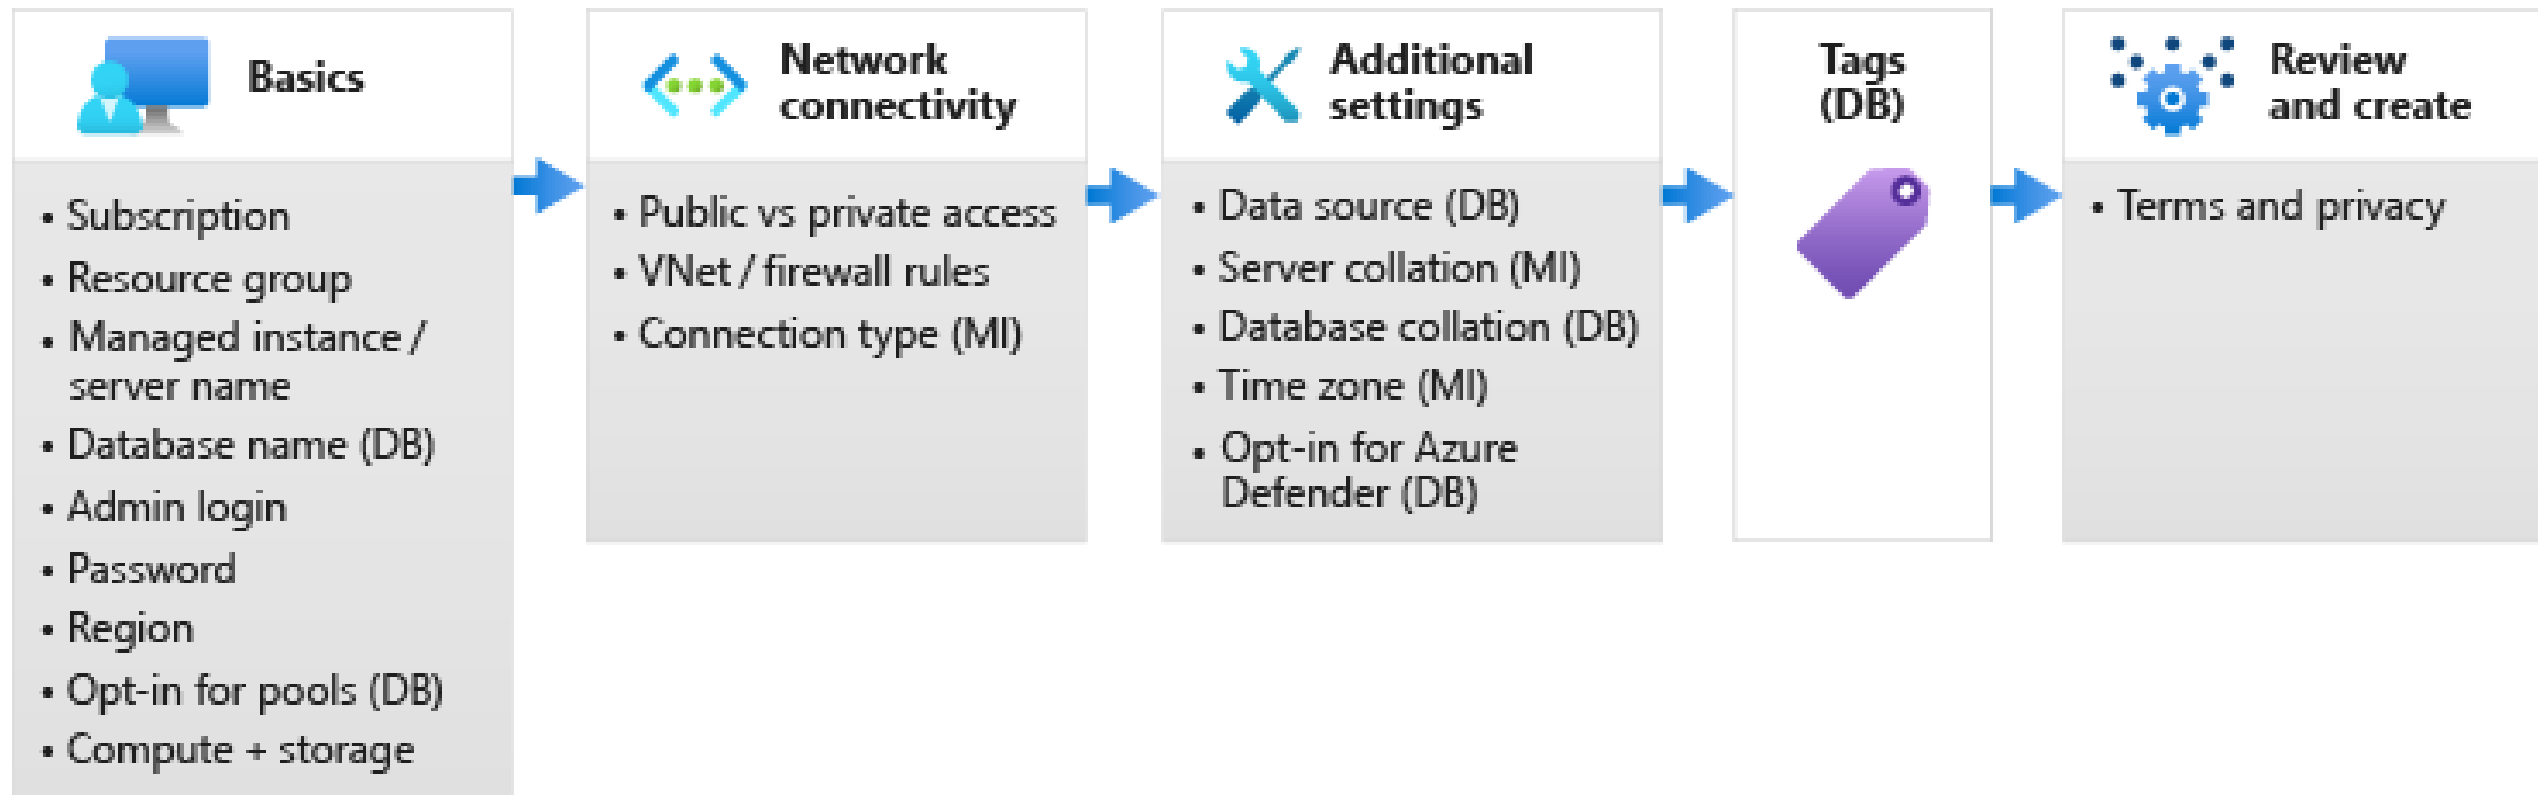
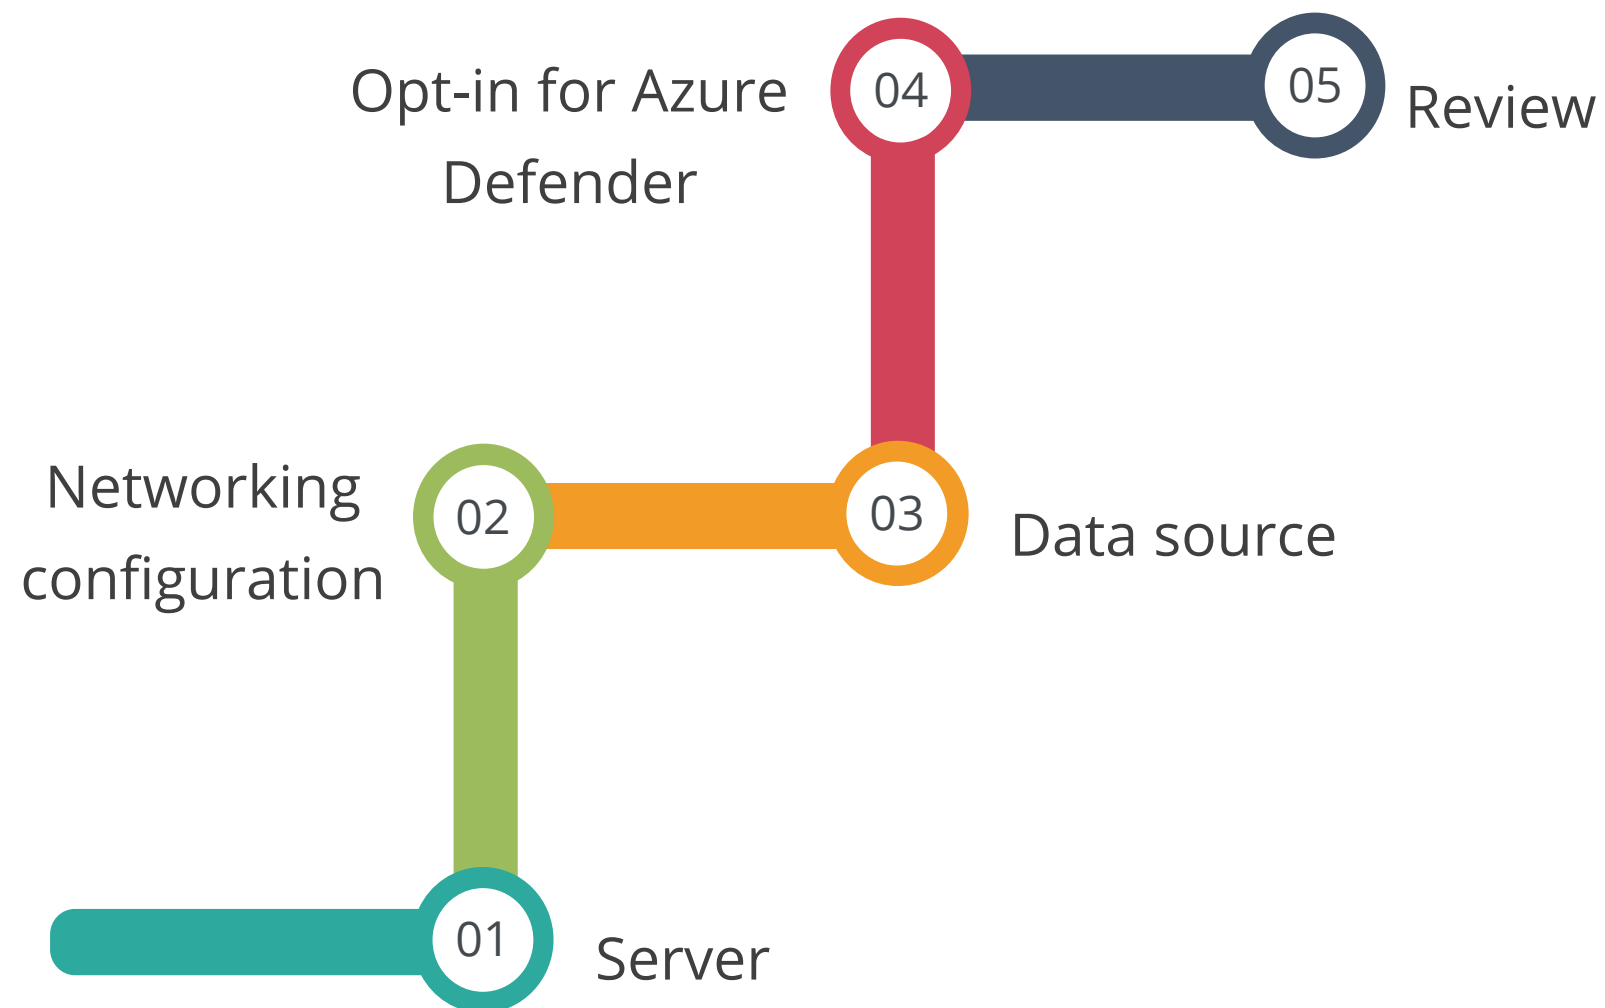


image source: <https://docs.microsoft.com/en-in/>

# Deploy Azure SQL Database

These are the five options from the sections in the Azure portal:



# Deploy Azure SQL Database

These are the explanations for the options:

## Networking Configuration

A public endpoint can be allowed to link to a controlled instance over the internet.

## Opt-in for Azure Defender

After the instance has been deployed, the user can allow Azure Defender.

## Server

The name of the server must be unique across Azure.

## Data Source

The user deploys the instance first and then the databases within it.

## Review


The user can check the deployment choices in the Review section.

# Recommending the Right Data Store

There are different options to choose from among SQL and NoSQL databases:

Relational database  
management systems

Key or value stores

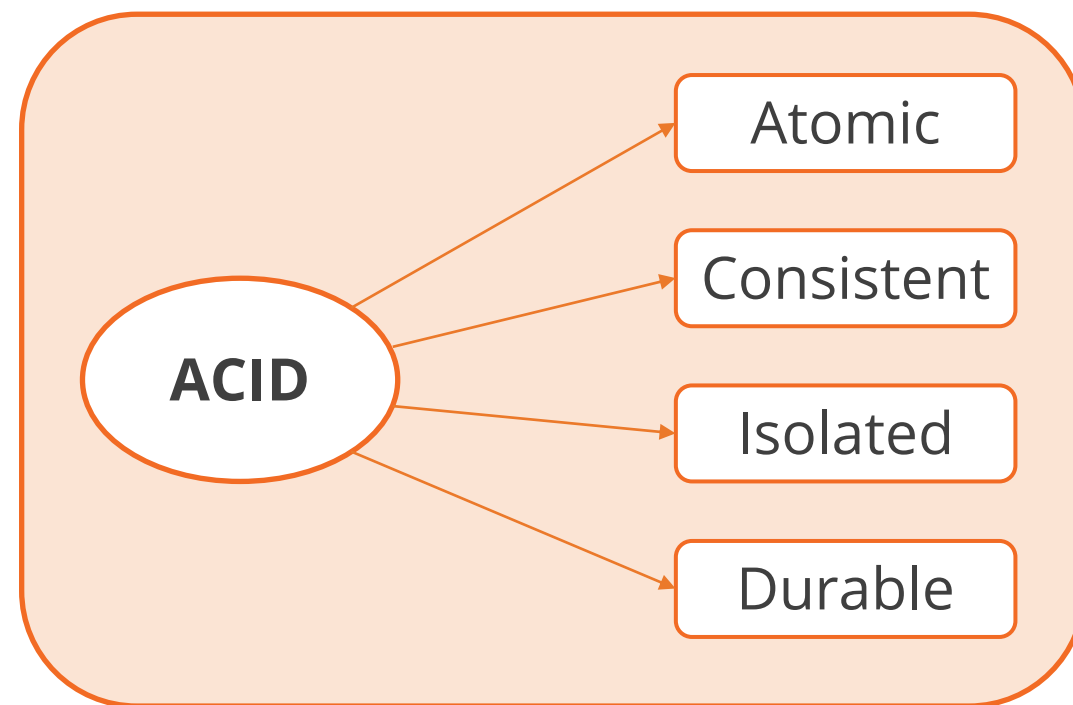


Graph databases

Document databases

# Relational Database Management Systems

Relational databases organize data as a series of two-dimensional tables with rows and columns.




Each table has its own columns, and every row in a table has the same set of columns.

RDBMS implements a transactionally consistent mechanism that conforms to the ACID model for updating information.

# Relational Database Management Systems

An RDBMS is useful when strong consistency guarantees are important.





# Key or Value Stores

A key or value store is a large hash table.


Key	Value
AAAAA	1101001111010100110101111...
AABAB	1001100001011001101011110....
DFA766	0000000000101010110101010...
FABCC4	1110110110101010100101101...

Opaque to  
data store

- It associates each data value with a unique key.
- Most key or value stores only support simple query, insert, and delete operations.
- An application can store arbitrary data as a set of values.
- Some key or value stores impose limits on the maximum size of values.

# Key or Value Stores

These are the features of key or value stores:



# Document Databases

It is conceptually like a key or value store, except that it stores a collection of named fields and data.

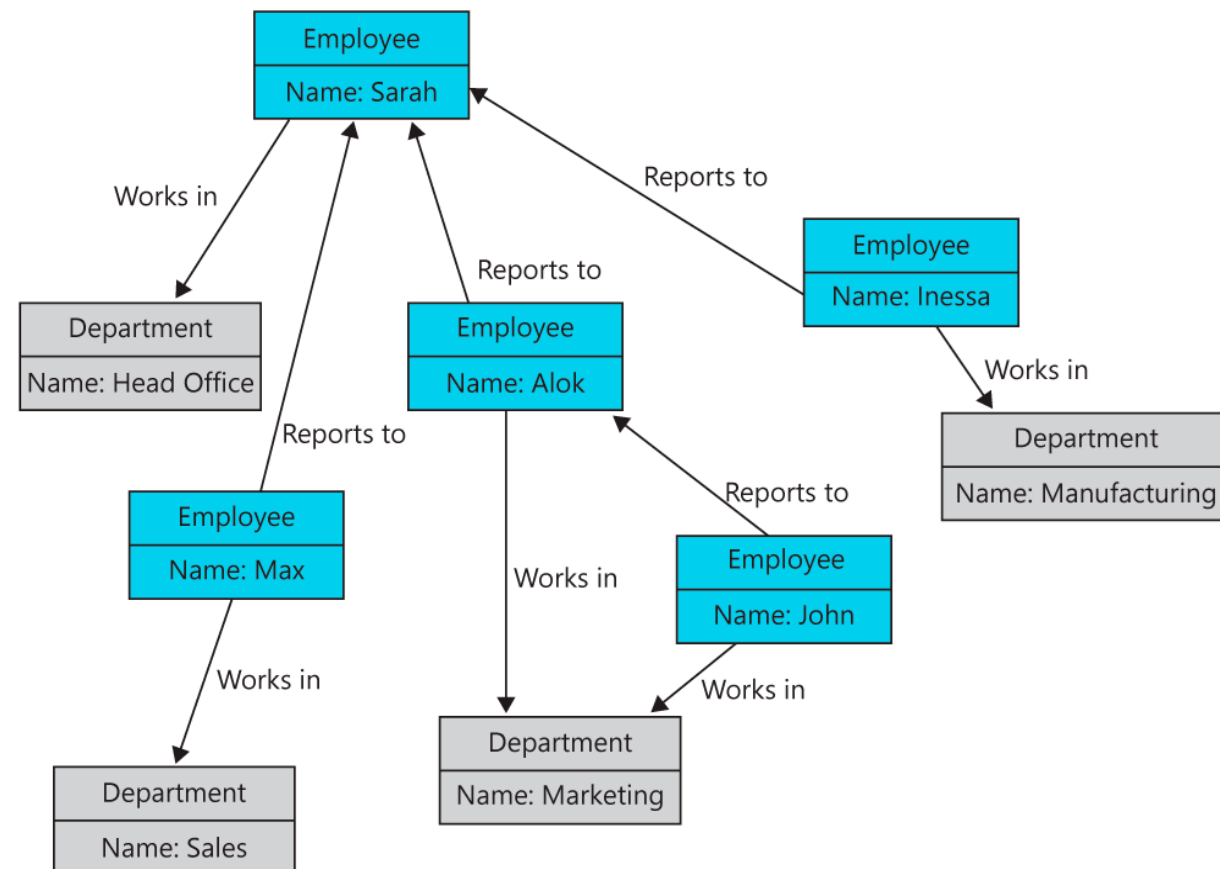
Data fields of a document can be encoded, including XML, YAML, JSON, and BSON.

Key	Document
1001	<pre>{   "CustomerID": 99   "OrderItems": [     { "ProductID": 2010,       "Quantity": 2,       "Cost": 520     },     { "ProductID": 4365,       "Quantity": 1,       "Cost": 18     }   ],   "OrderDate": "04/01/2017" }</pre>
1002	<pre>{   "CustomerID": 220   "OrderItems": [     { "ProductID": 1825,       "Quantity": 1,       "Cost": 120     }   ],   "OrderDate": "05/08/2017" }</pre>

- A document store does not require that all documents have the same structure.
- **Relevant Azure services:** Cosmos DB

# Graph Databases


A graph database stores two types of information: nodes and edges.



- Nodes store data entities.
- Edges specify the relationships between nodes.
- Both nodes and edges can have properties that provide information about a particular node.
- Edges can also have a direction indicating the nature of the relationship.
- **Relevant Azure services:** Cosmos DB

# Select a Candidate Data Store

Azure offers several managed data storage solutions; each one provides different features and capabilities.

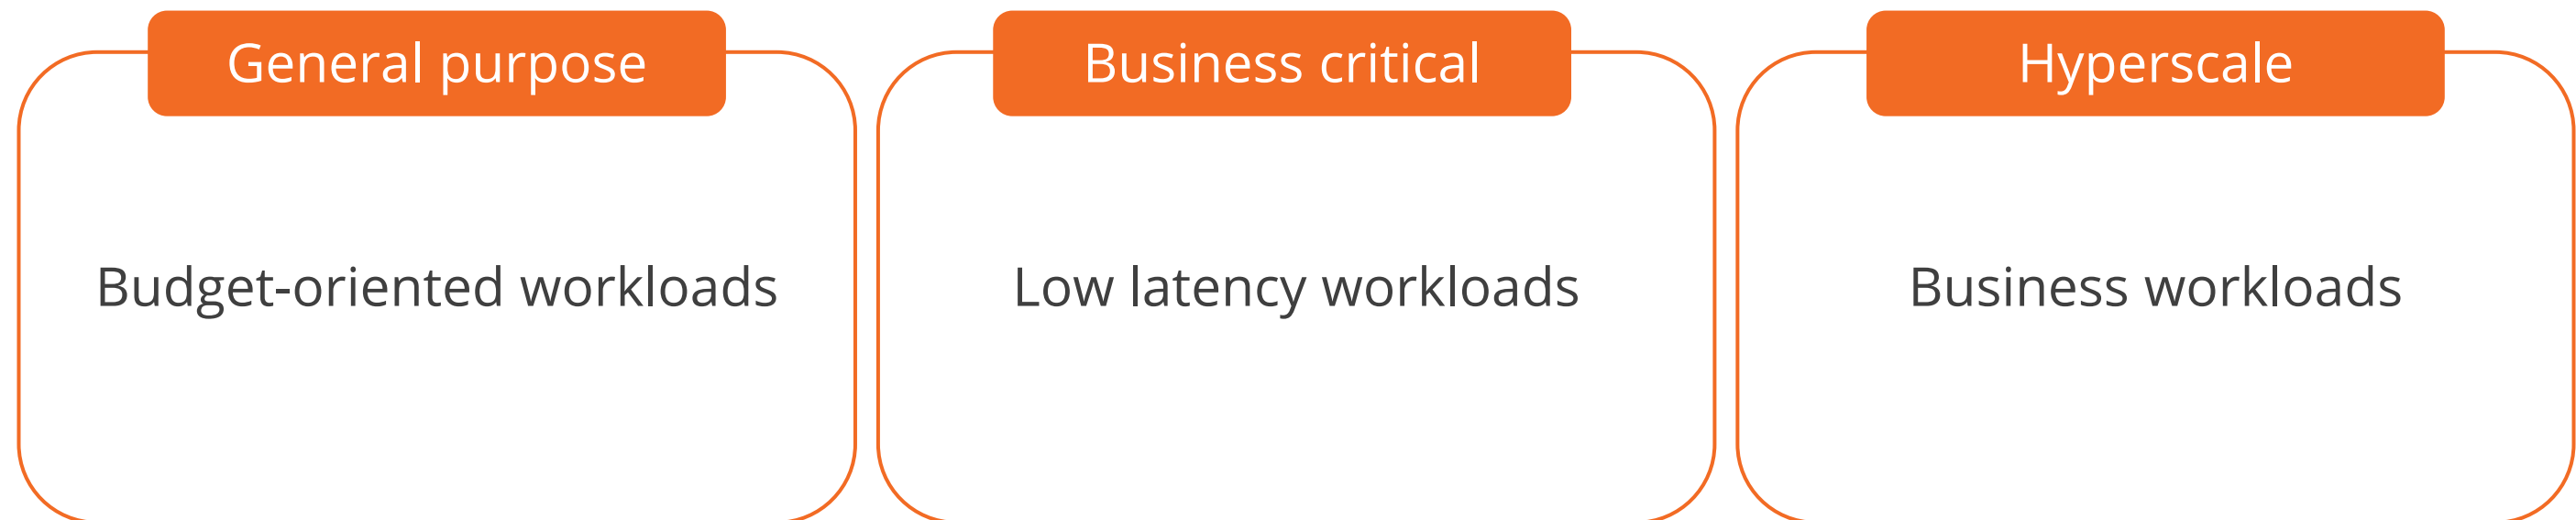


- If the application consists of multiple workloads, evaluate each workload separately.
- A complete solution may incorporate multiple data stores.

# Recommend Database Service Tier Sizing

# Azure SQL Database and Azure SQL Managed Instance

The types of Azure SQL database & SQL managed instance service tiers are:



# General Purpose Service Tier

This service tier comes with 99% availability.

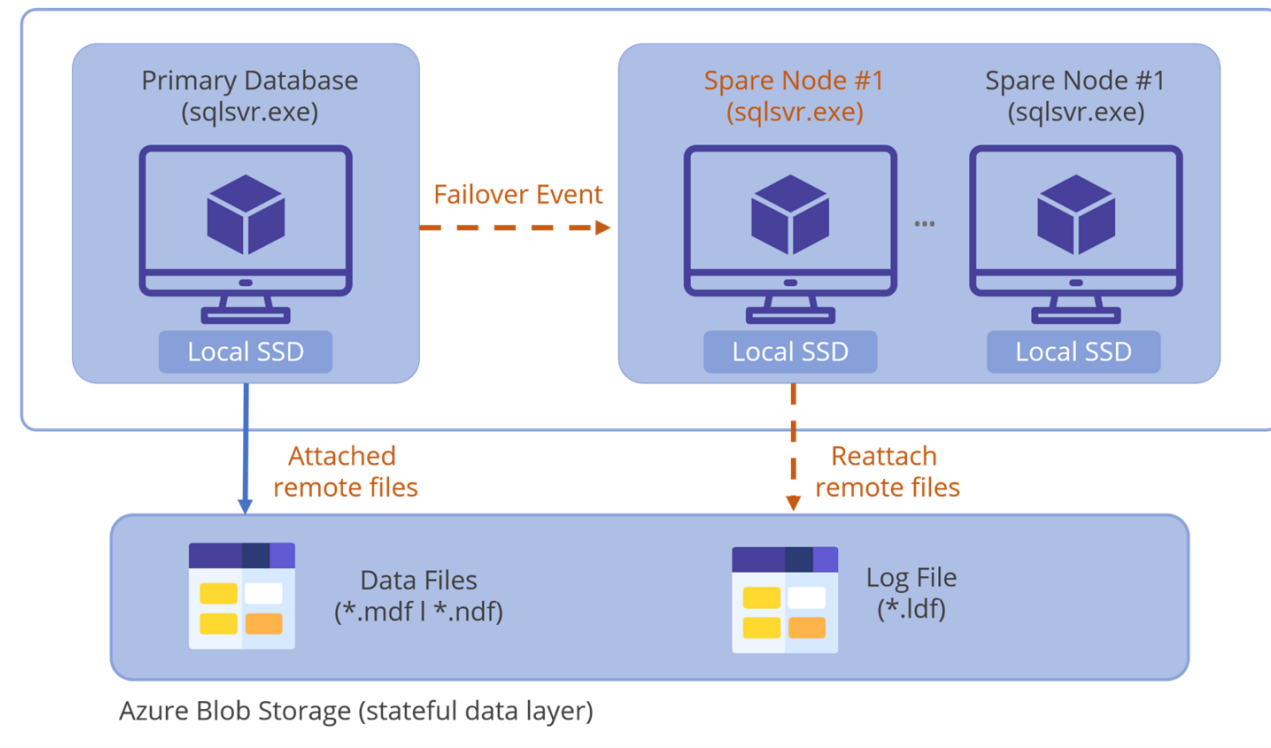
## System Diagram

General Purpose – Azure SQL database

Azure Data Centre

(**userdb**).database.windows.net

Azure App Fabric (stateless computing layer)

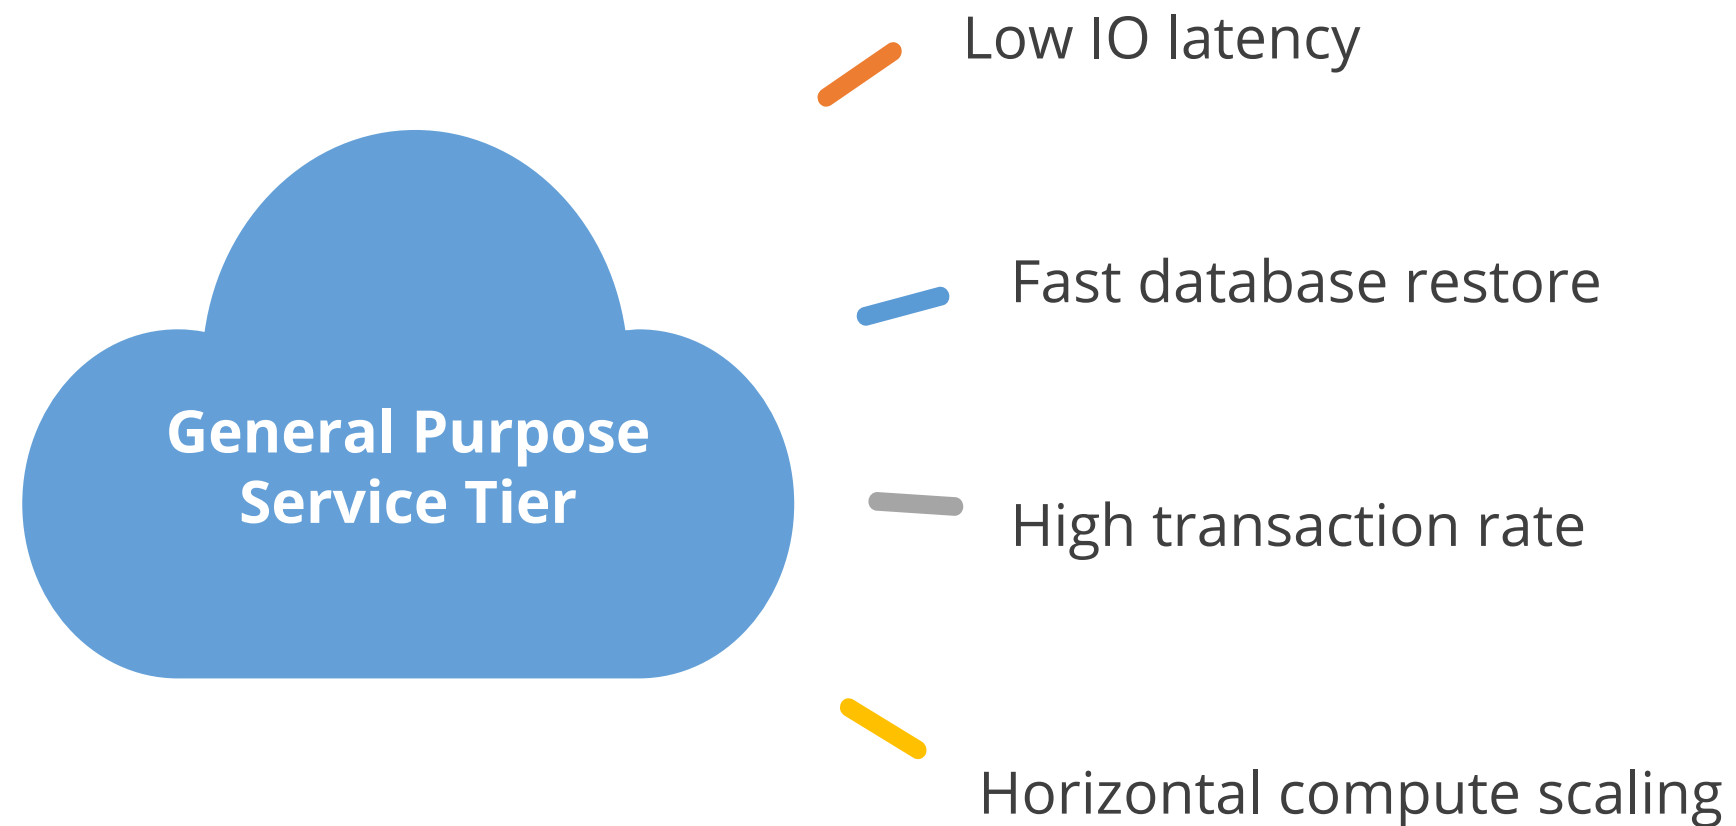


It is a default service tier that is suitable for most generic workloads.



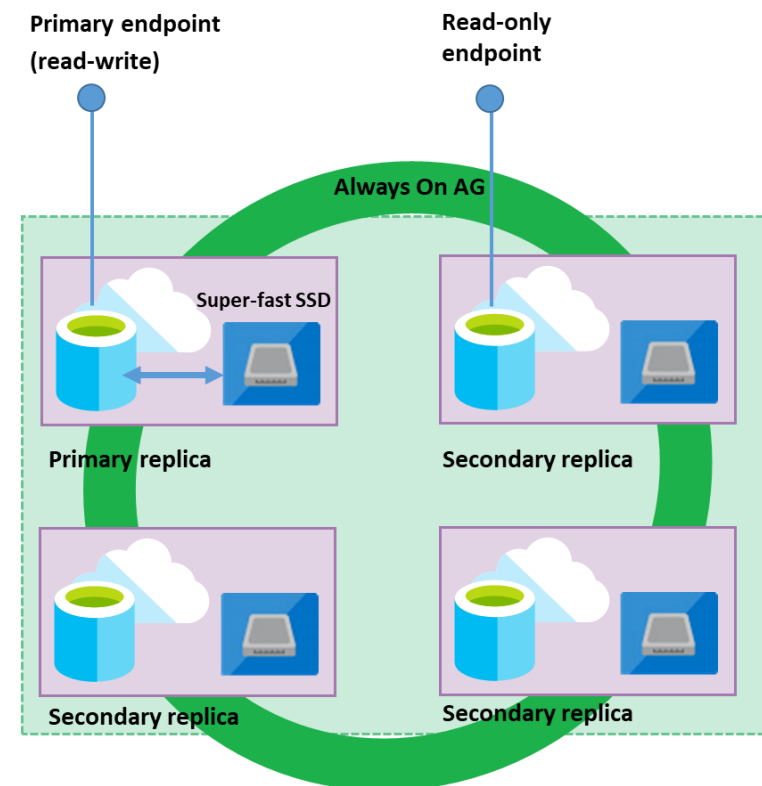
# General Purpose Service Tier

The benefits of General Purpose service tier are:



# Business Critical Service Tier

The Business Critical Service Tier is based on a cluster of database engine processes.




Business Critical service tier: collocated compute and storage

Premium availability is enabled in this service tier.

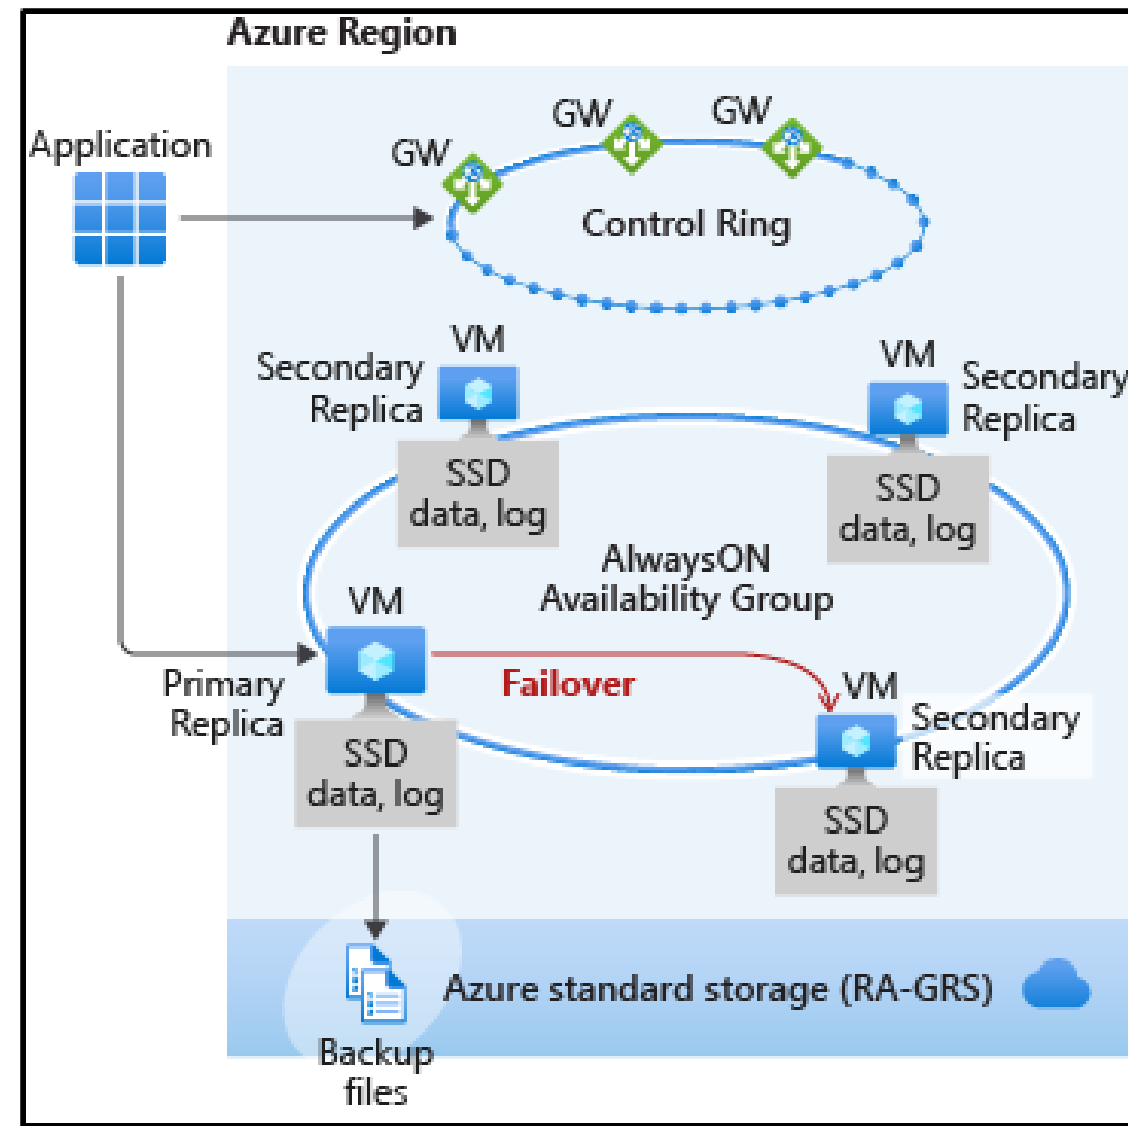
# Business Critical Service Tier

Business critical service tier can be chosen when there are:



# Business Critical Service Tier Availability

Business critical high availability service tier is created when a four-node cluster is formed.



# Service Tier Comparison

The difference between the service tiers is given in the table below:

Availability	Resource type	General Purpose	Business Critical	Hyperscale
Compute size	SQL database	1 to 80 vCores	1 to 80 vCores	1 to 80 vCores
Storage size	SQL database	5 GB - 4 TB	5 GB - 4 TB	Upto 100 TB
TempDB size	SQL managed instance	24 GB per vCore	Upto 4 TB	N/A
In-memory OLTP	SQL managed instance	N/A	Available	N/A
Database size	SQL database	5 GB - 4 TB	Upto 100 TB	5 GB - 4 TB

# Service Tier Comparison

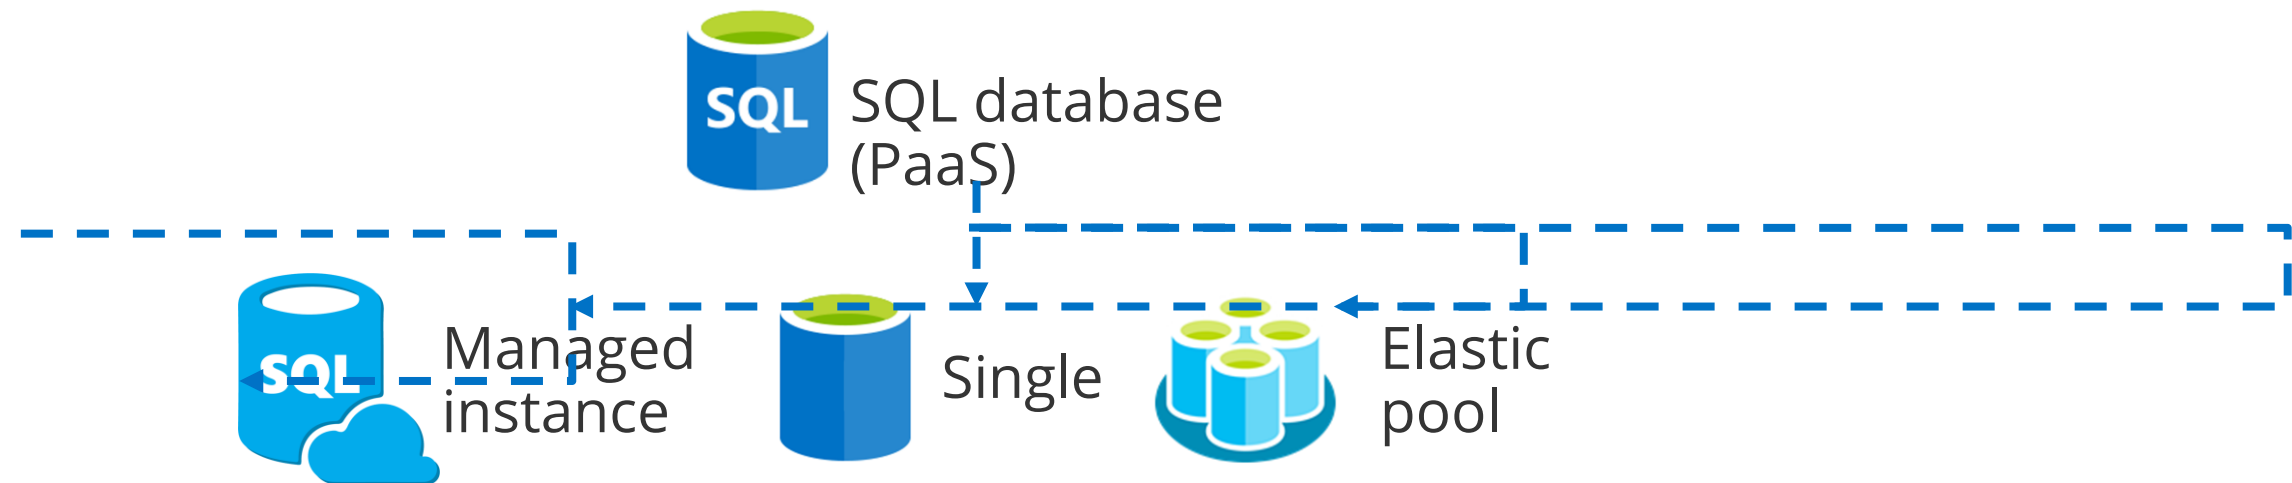
General Purpose	Hyperscale	Business Critical
Least expensive	Less expensive	Most expensive
Latency 5-10 milliseconds	Scales compute resources up and down very quickly	Latency 1-2 milliseconds
99.99 % availability	Instant backups and fast database restores	99.95 % availability
Maximum size 4TB	Maximum size 100TB	Local SSDs on four-node cluster

The service tier characteristics might be different in SQL database and SQL managed instance.

## Recommend a Solution for Database Scalability

# Dynamically Scale Azure SQL Database and Managed Instance

SQL database enables the users to change resources allocated to the databases.



SQL managed instance is a new deployment option that enables frictionless migration for SQL apps and modernization in a fully managed service.



# Dynamically Scale Azure SQL Database and Managed Instance

These are the benefits of Scaling:




- It is cost effective.
- It is more versatile.
- It mitigates performance issues.
- It handles the incoming workload well.

# Azure SQL Database and Managed Instance Pricing

Azure SQL database and SQL managed instance offers two types of purchasing model:

DTU-based  
purchasing model




vCore-based  
purchasing model

**Purchasing model**

# The DTU-Based Purchasing Model

A Database Transaction Unit (DTU) is a composite metric that includes CPU, memory, reads, and writes.



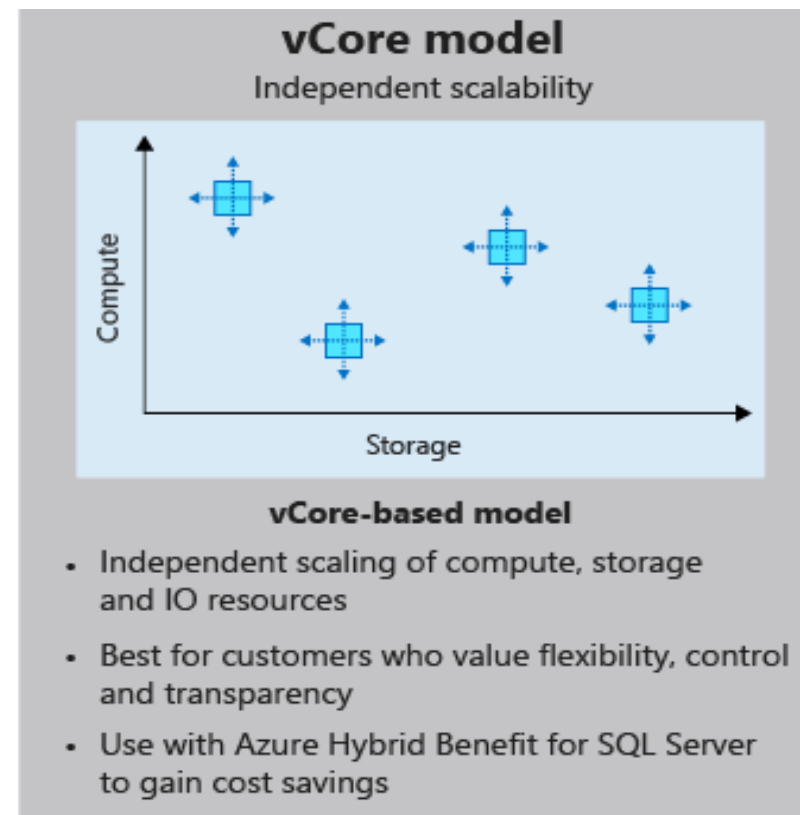
## Database Transaction Unit (DTU)-based model

- Bundled measure of compute, storage and IO resources
- Best for customers who want simple, pre-configured resource options

The DTU-based purchase model provides a list of preconfigured compute resource bundles.

# The vCore-Based Purchasing Model

The vCore-based model helps in optimizing price.



The vCore-based pricing strategy allows users to choose computing and storage resources individually.


# Scaling Azure SQL Database and Managed Instance

Scaling the database can be done via the Azure portal using a slider.



# Scale Single Databases in Azure SQL Database

Azure SQL Database allows the users to scale the databases dynamically.



# Scale Single Databases in Azure SQL Database

The user can determine the maximum number of resources that will be assigned to each database using either of the purchasing models.



The user can set a maximum resource limit for each group of databases in an elastic pool.

# Transport Layer Security

Transport Layer Security(TLS) encrypts communication between a client application and an Azure Storage account.



TLS is a standard cryptographic protocol that assures client and service privacy and data integrity over the Internet.



# Transparent Data Encryption

SQL Server, Azure SQL Database, and Azure Synapse Analytics data files are encrypted using transparent data encryption (TDE).



Encrypting data at rest is the terminology for this type of encryption.

# Transparent Data Encryption

Precautions can be taken to secure a database, such as:



Creating a safe system



Encrypting sensitive information



Constructing a firewall to  
protect the database  
servers

# Azure Security Audit

The audit methodology is based on Microsoft's own best practice recommendations for Azure, and it primarily covers the following topics:



- Improve user's Microsoft Defender for Cloud security score
- Use an industry accepted benchmark to assess an organization's current security posture

# Azure Security Audit



- Conduct internal and external compliance audits on a regular basis, as well as regulatory compliance attestations
- Examine the policy's specifications
- For a comprehensive perspective of the technical Azure Governance setup, use Azure Governance Visualizer

# Assisted Practice

## Azure SQL DB

Duration: 10 Min.

### Problem Statement:

You've have been assigned a project to construct an Azure SQL DB that will automate most database maintenance tasks such as updating, patching, backups and monitoring without user involvement.

# Assisted Practice: Guidelines

Steps to create an Azure SQL DB are:

1. Login to your Azure portal
2. Search and select Azure SQL
3. Select the SQL database tile
4. Fill in the required fields and create SQL DB



# Design Security for Data


# Design Security For Data

Data on the Azure platform is always encrypted, and Azure provides customers with strong data security, both by default and as customer options.


Data at Rest



Data in Transit



Data in Use





# Data at Rest


Data at rest is static data stored on hard drives that are archived or not frequently accessed or modified, and the protection at rest aims to secure the inactive data stored on devices or networks.



- Available for the services across the software as a service (SaaS)
- Available for the platform as a service (PaaS)
- Available for the infrastructure as a service (IaaS) cloud models

# Data in Transit

Data in transit is the data actively moving from one location to another, such as across the internet or through a private network.



- Data loss prevention (DLP) technology
- Backup and recovery solutions
- Identity and access management technology

## Data in Use

Data in use is to protect the confidentiality and integrity of data while it is in use, in addition to existing protection to encrypt data in transit and at rest.



Available on a new DC-series of virtual machines

# Design for Azure SQL Edge

# Azure SQL Edge

Azure SQL Edge is an optimized relational database engine geared for IoT and IoT Edge deployments.



- It provides capabilities to create a high-performance data storage and processing layer for IoT applications.
- Azure SQL Edge can stream, process, and analyze relational and non-relational data such as JSON, graph, and time-series data.

# Azure SQL Edge

Azure SQL Edge provides the same Transact-SQL (T-SQL) programming surface area that makes the development of applications or solutions easier and portability among IoT Edge devices, data centers, and the cloud straightforward.



# Advanced Security of Azure SQL Edge

The features that provide advanced security of Azure SQL Edges are:



Transparent data encryption and Always Encrypted features that provide data security



Role-based access control (RBAC) and attribute-based access control (ABAC)



Compliance with security rules with the help of data classification

# Supported Features of Azure SQL Edge

The supported features of Azure SQL Edge are:

1

SQL streaming based on the same engine that powers Azure Stream Analytics and provides real-time data streaming capabilities in Azure SQL Edge

2

The T-SQL function call **Date\_Bucket** for Time-Series data analytics

3

Machine learning capabilities through the ONNX runtime, included with the SQL engine



# Unsupported Features of Azure SQL Edge

The Unsupported features of Azure SQL Edge are:

## Database Design

Full-text indexes and search, related DDL commands and Transact-SQL functions, catalog views, and dynamic management views

## Database Engine

CLR assemblies, related DDL commands and Transact-SQL functions, catalog views, and dynamic management views

# Un-Supported Features of Azure SQL Edge

Subsystems:  
CmdExec, PowerShell,  
Queue Reader, SSIS,  
SSAS, and SSRS


SQL Server Agent

SQL Server Utility  
Control Point


Manageability

# Azure SQL Edge Editions

Azure SQL Edge is available with two different editions or software plans such as:




# Azure SQL Edge Working



Azure SQL Edge makes it easier and more productive to design and maintain applications. Users can build fantastic apps and solutions for their IoT Edge needs using all of their familiar tools and abilities.

# Azure SQL Edge Working

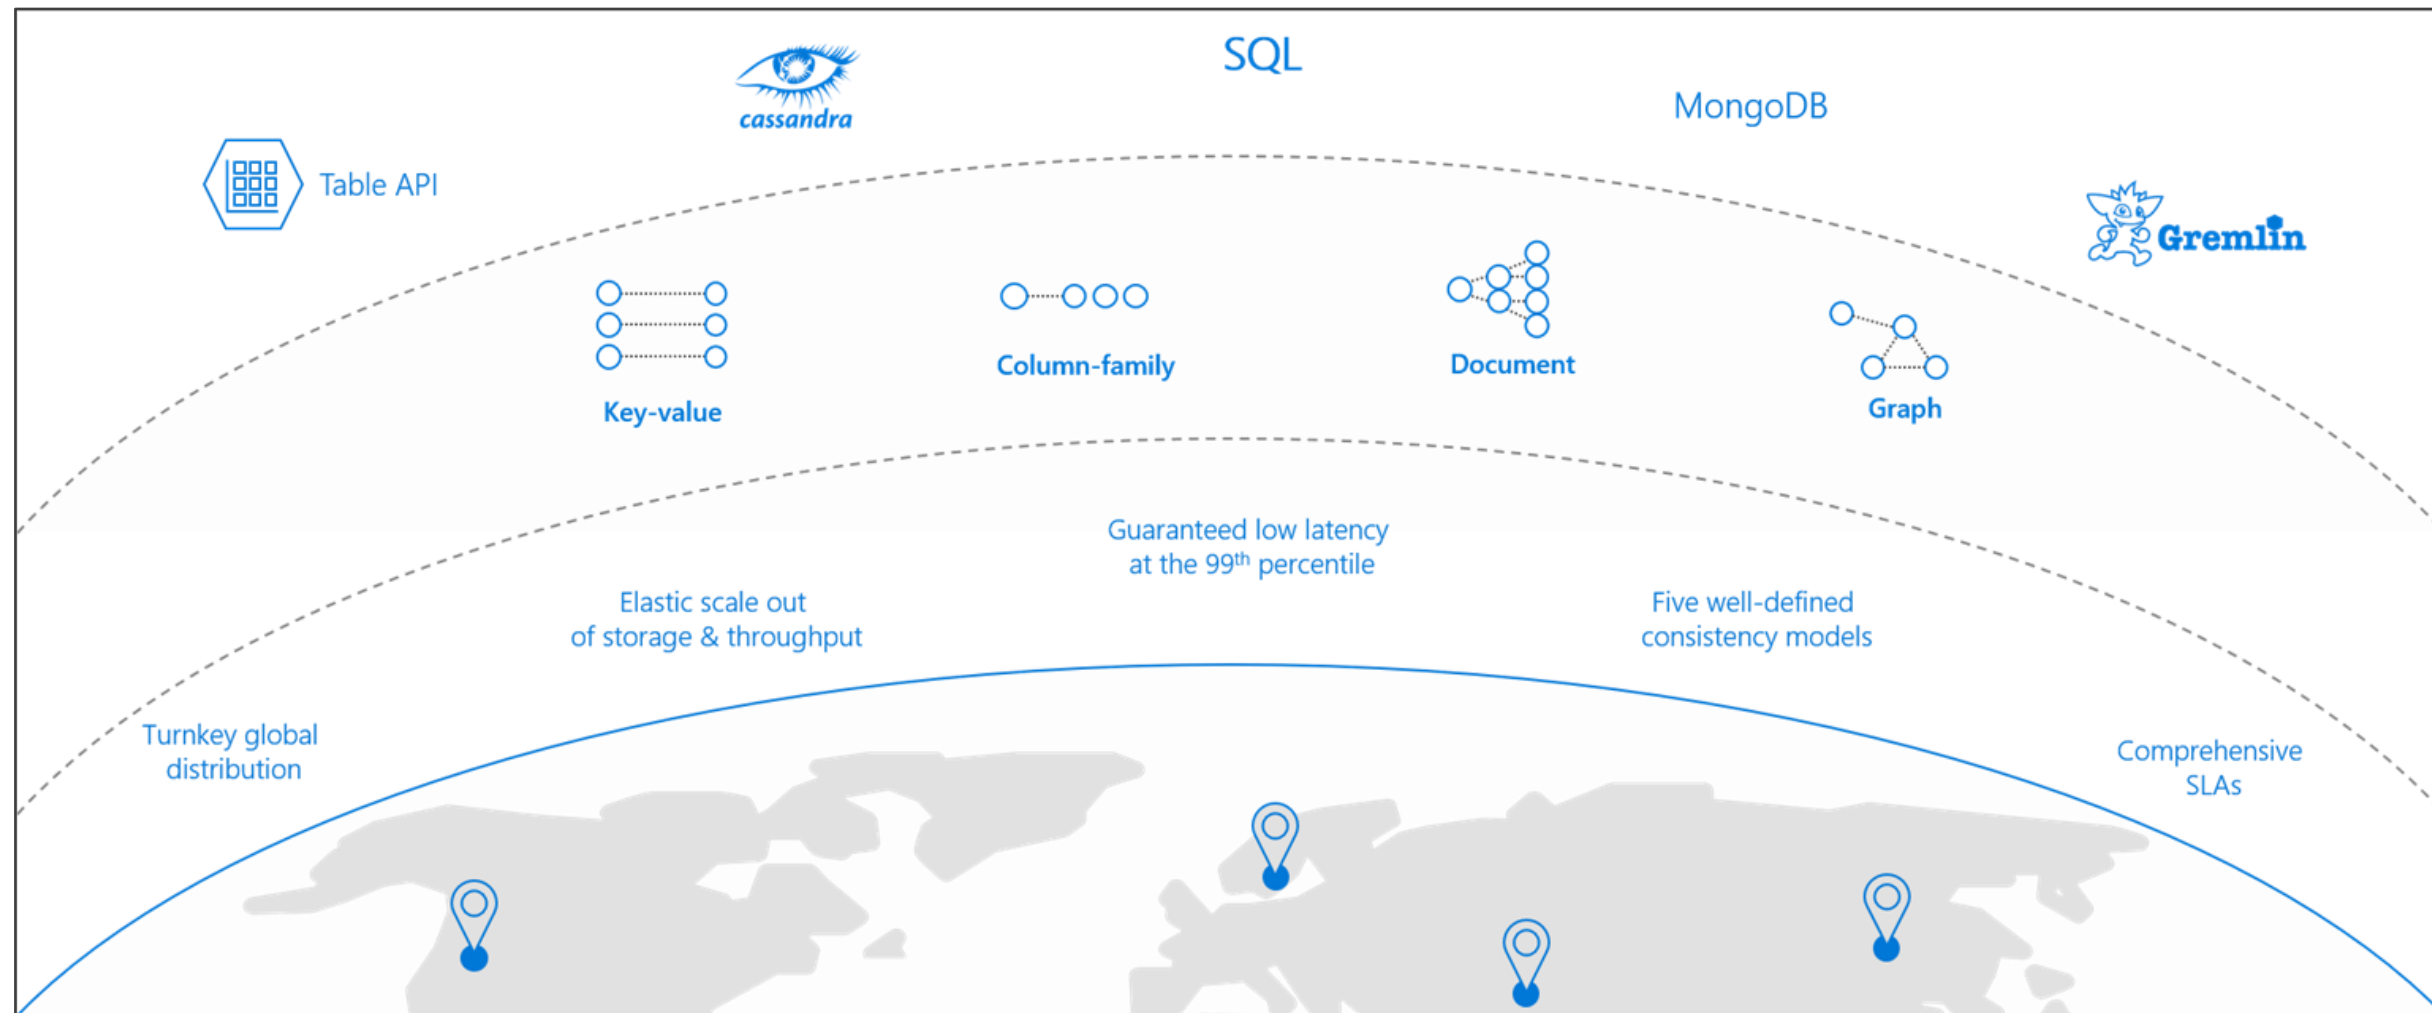
A user can develop in SQL Edge using tools such as:



# Design for Azure Cosmos DB

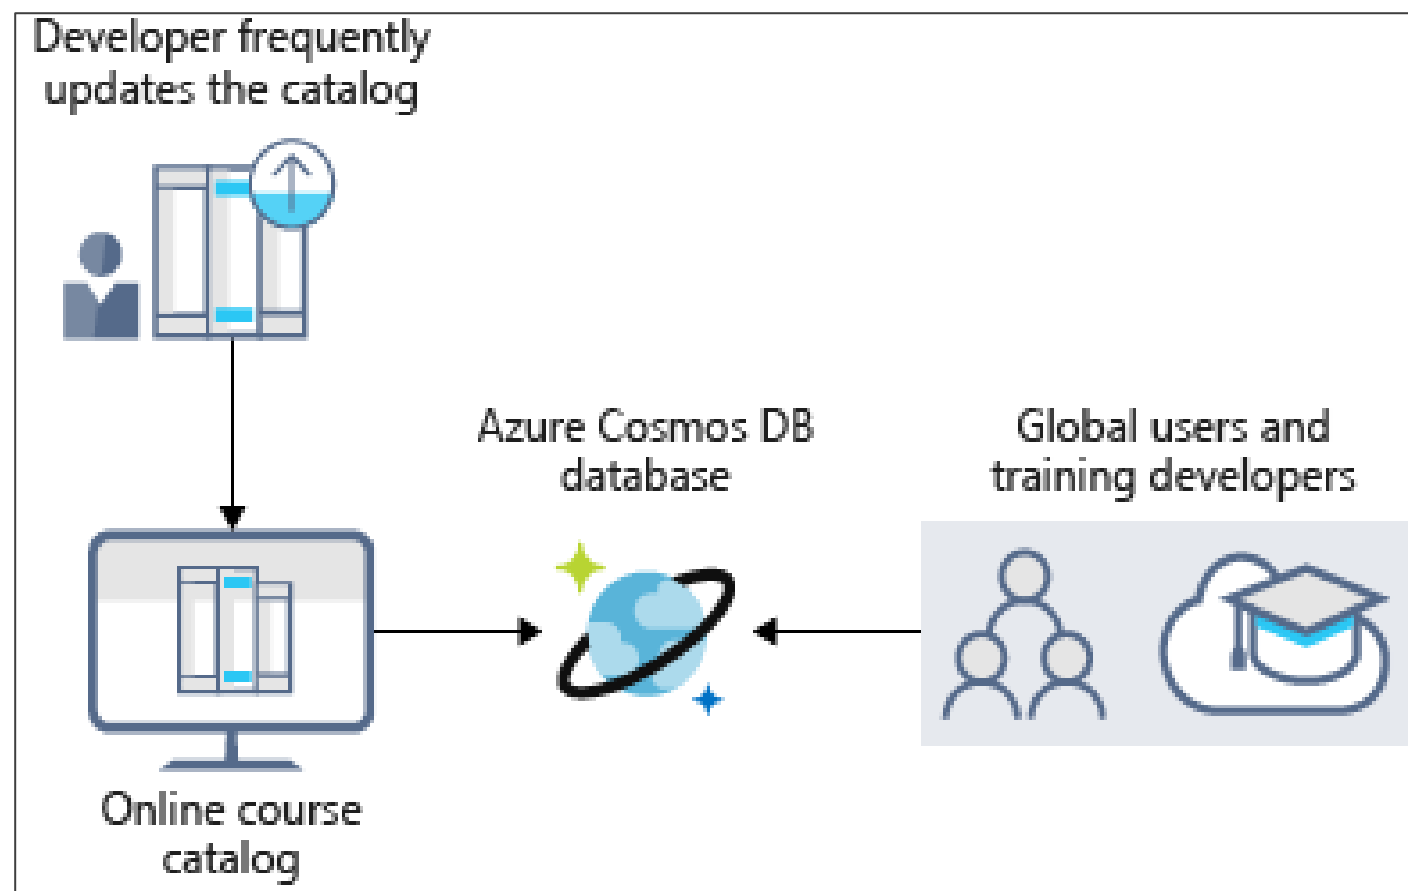
# Azure Cosmos DB

Azure Cosmos DB is a globally distributed and elastically scalable database.



# Azure Cosmos DB

The key benefits of Azure Cosmos DB are:



- Geo-replication
- Elastic scaling of throughput and storage worldwide
- Five well-defined consistency levels



# Database API in Azure Cosmos DB

Decision tree to choose an API in Azure Cosmos DB :

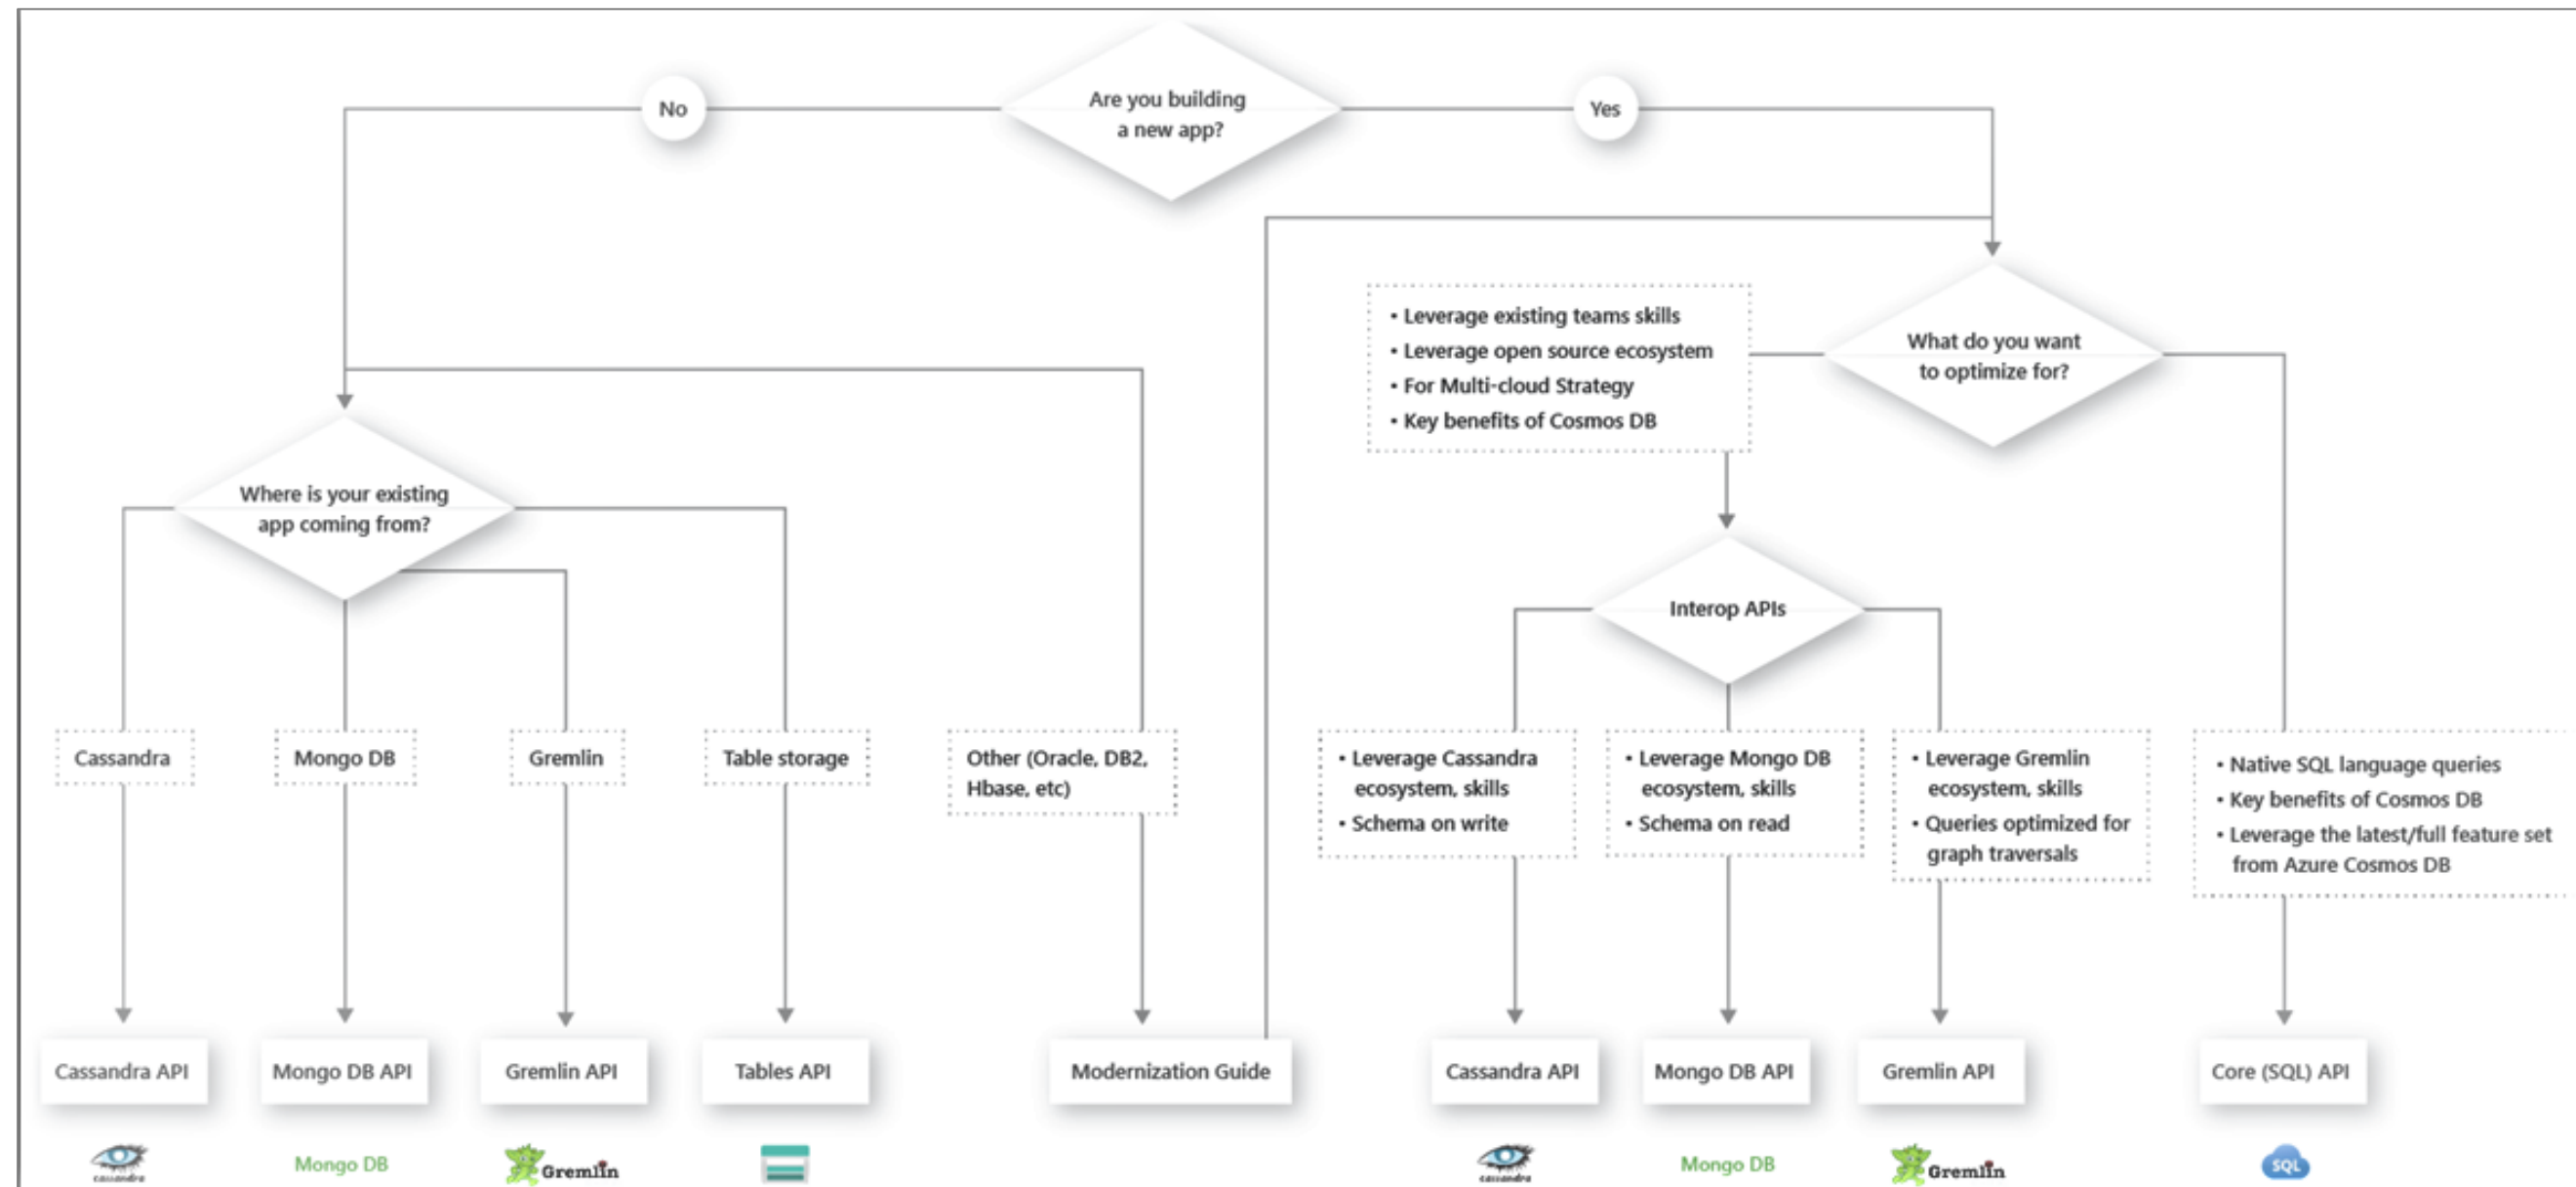



Image Source: <https://docs.microsoft.com/>

# Database API in Azure Cosmos DB

Azure Cosmos DB offers multiple database APIs, such as:



Core (SQL) API

API for MongoDB

Cassandra API

Gremlin API

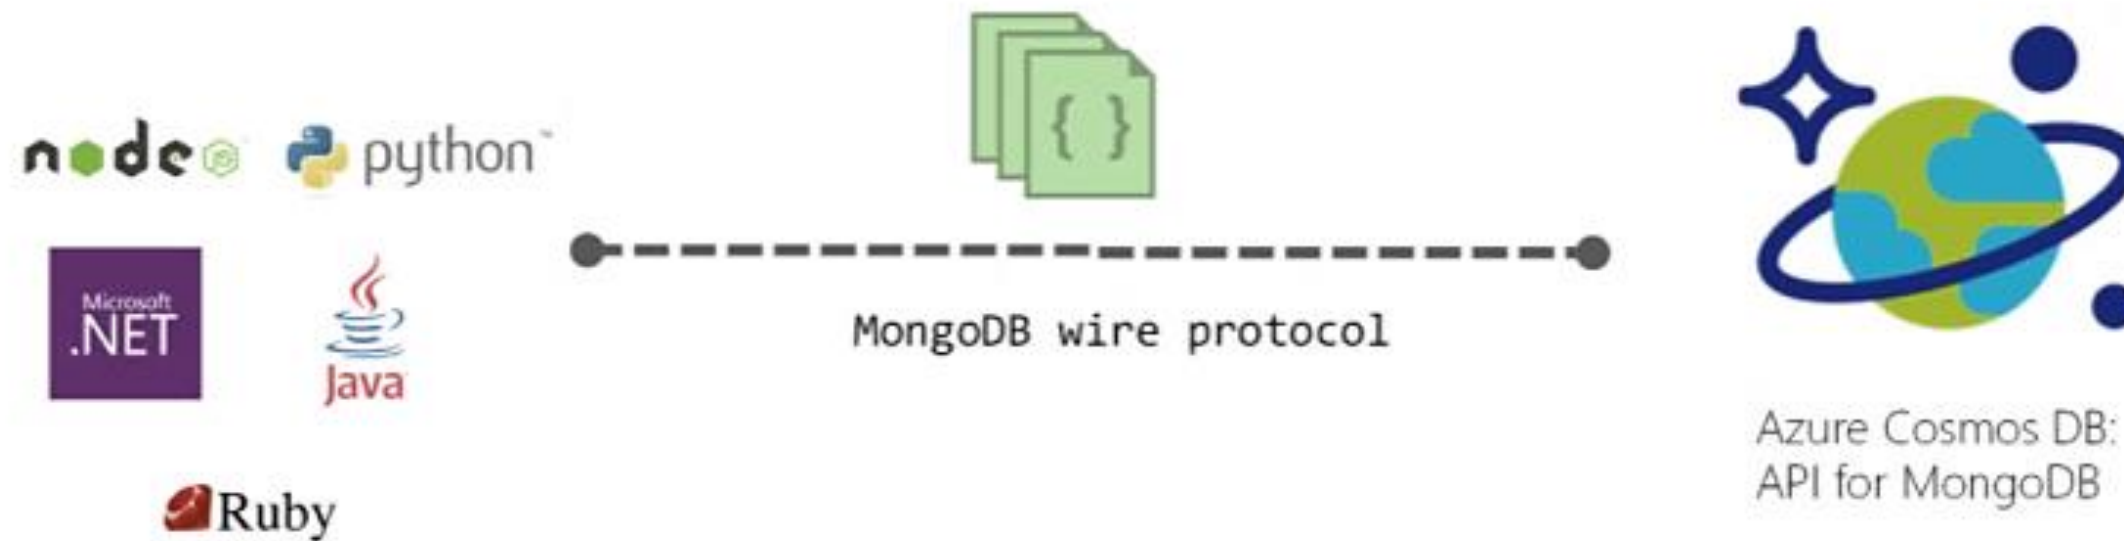
## Core (SQL) API

The feature that is being rolled out to Azure Cosmos DB is first available on SQL API accounts.



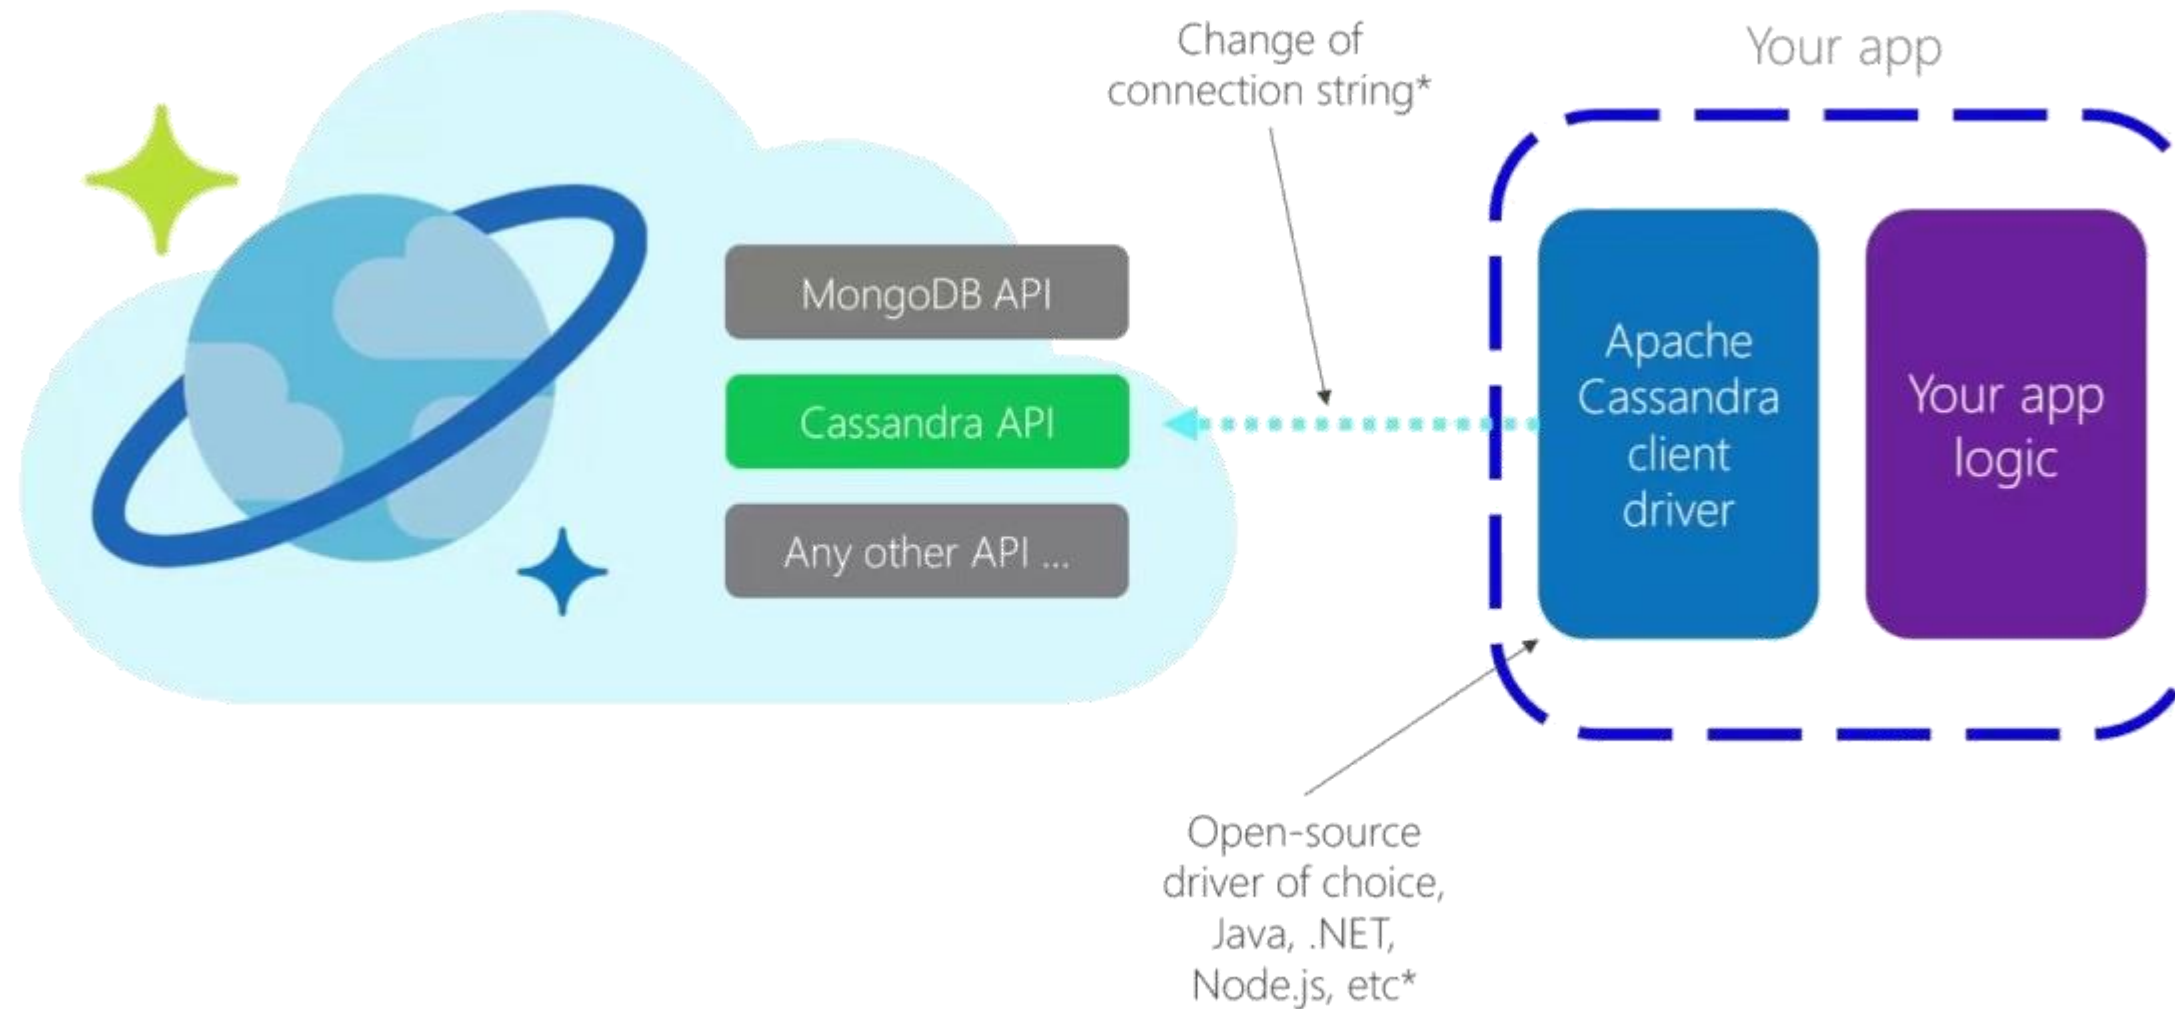
# API for MongoDB

The API for MongoDB enables the use of existing client drivers by adhering to the MongoDB wire protocol.



# Cassandra API

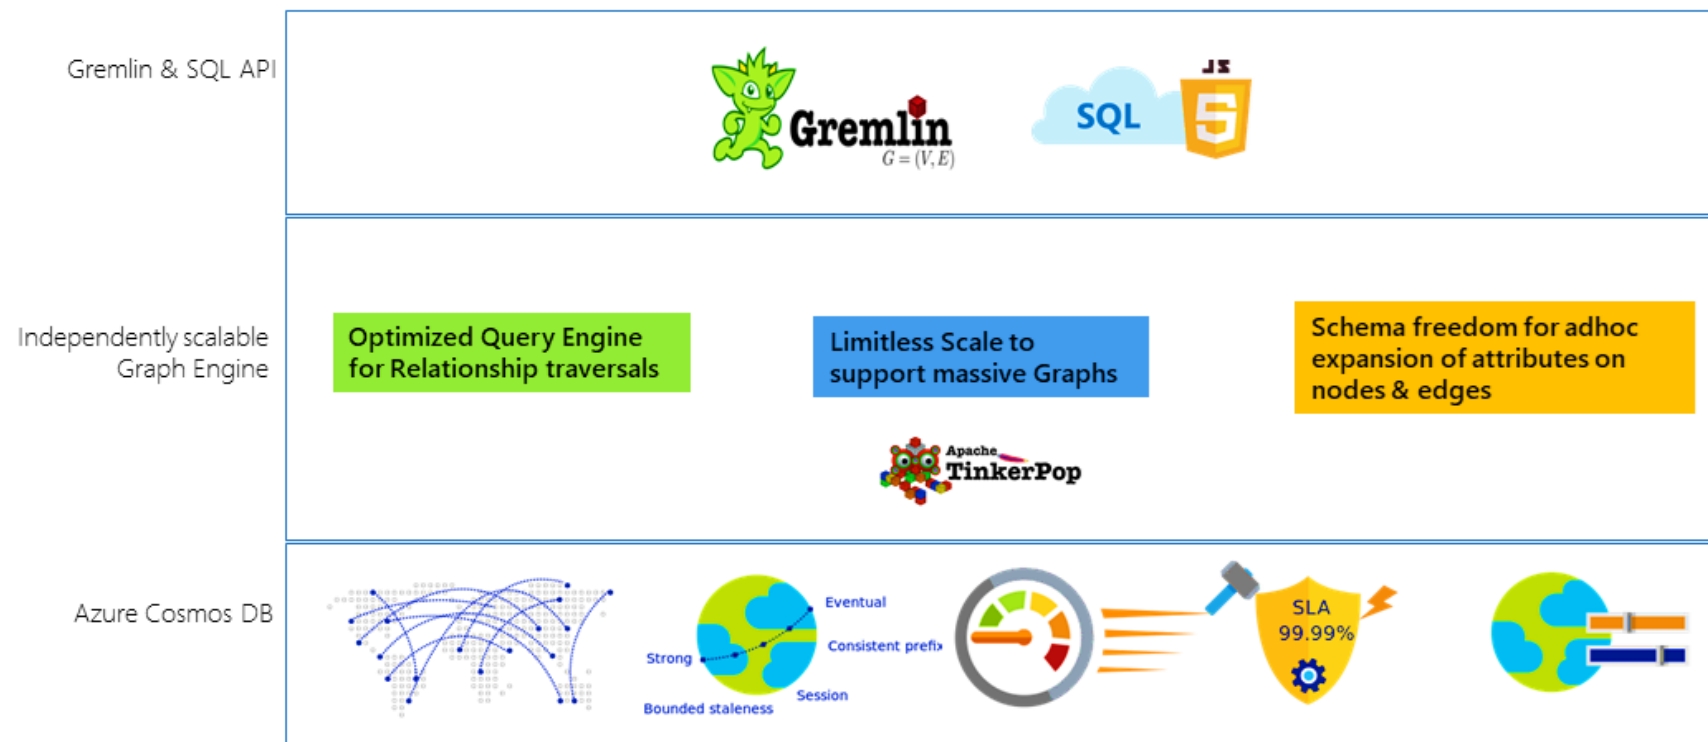
Cassandra API in Azure Cosmos DB aligns with the philosophy of approaching distributed NoSQL databases, and the API stores data in a column-oriented schema.



# Gremlin API


Azure Cosmos DB's Gremlin API combines the power of graph database algorithms with highly scalable, managed infrastructure.

## Azure Cosmos DB – Graph API PaaS



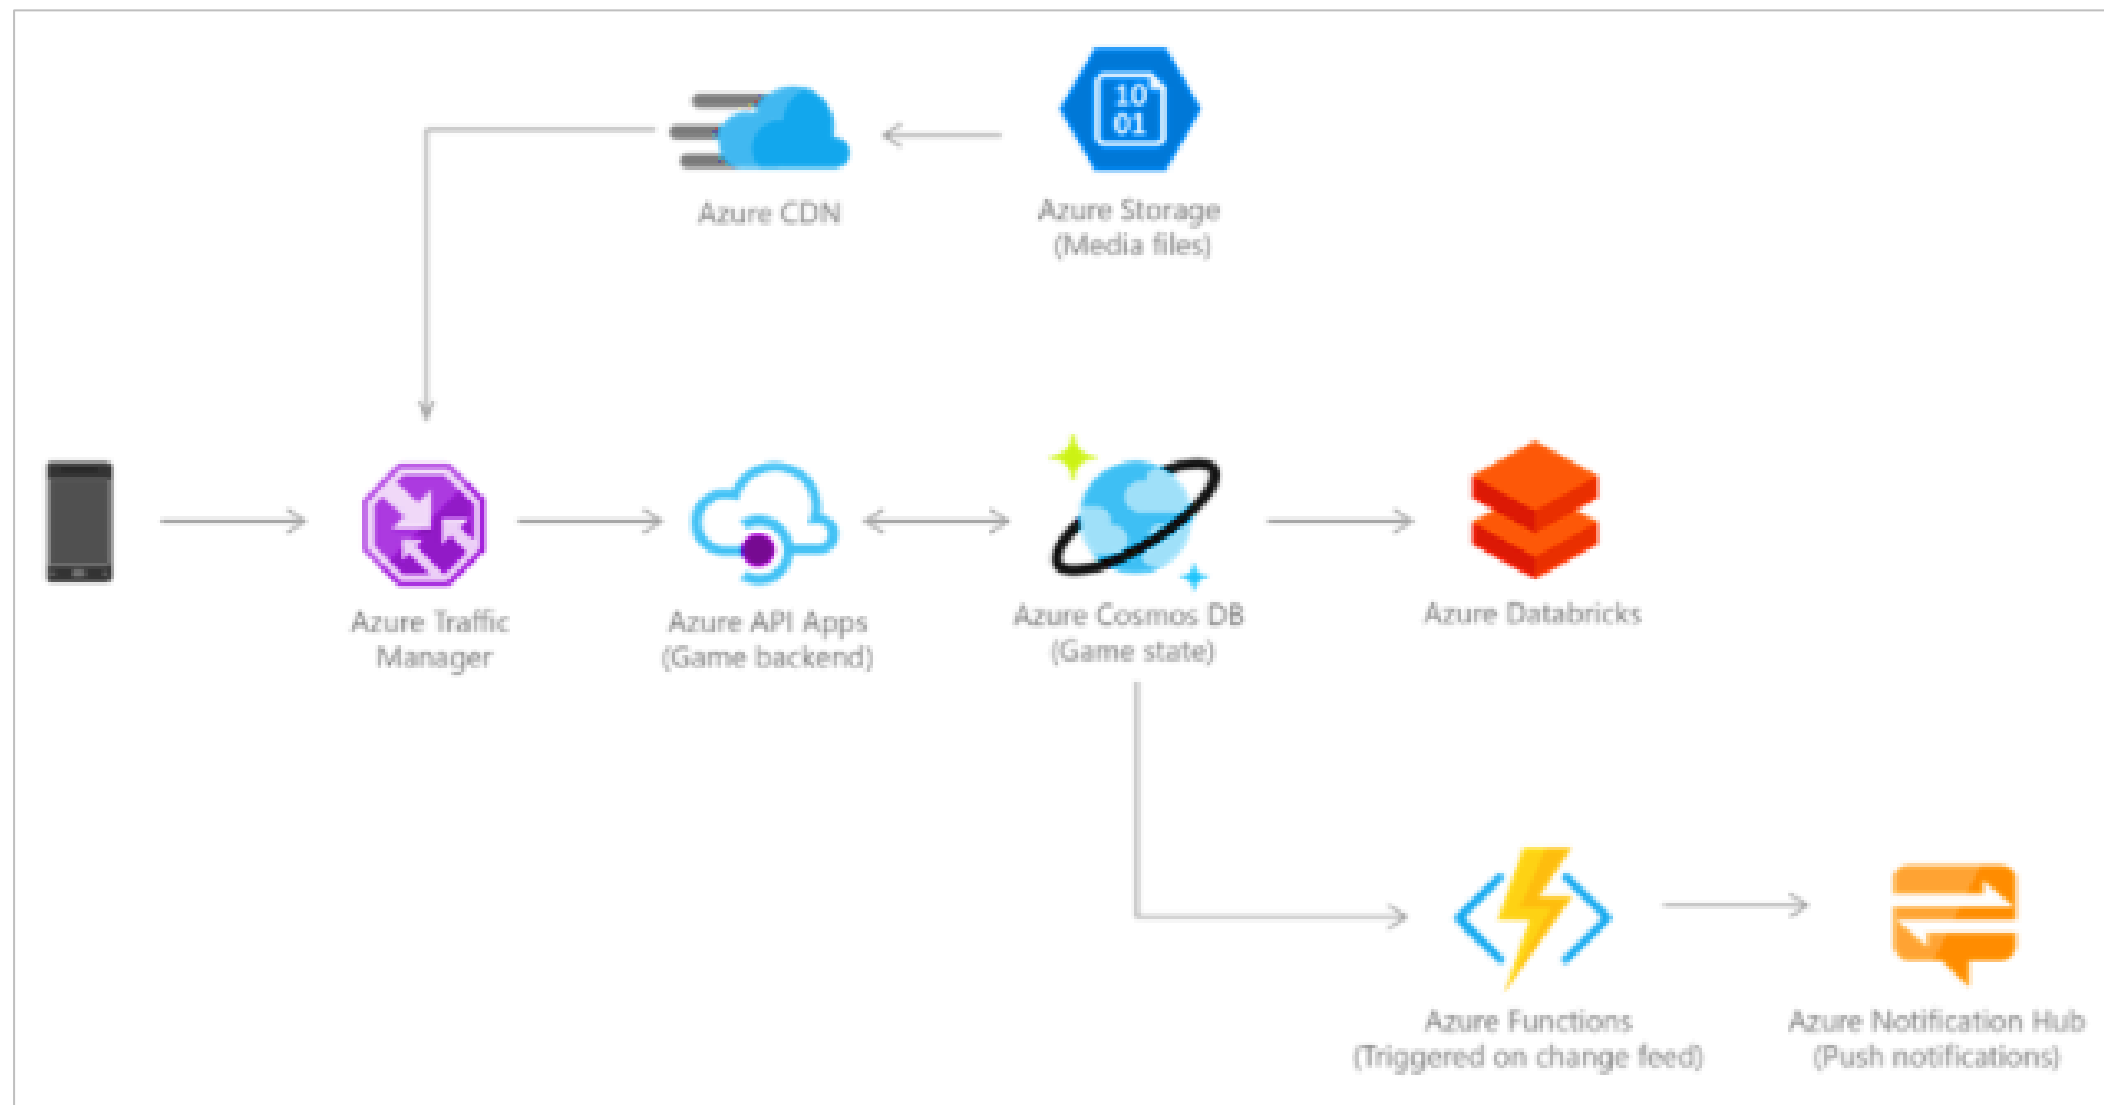
# Use Case of Azure Cosmos DB

Applications that integrate with third-party social networks must respond to changing schemas from these networks.



# Use Case of Azure Cosmos DB

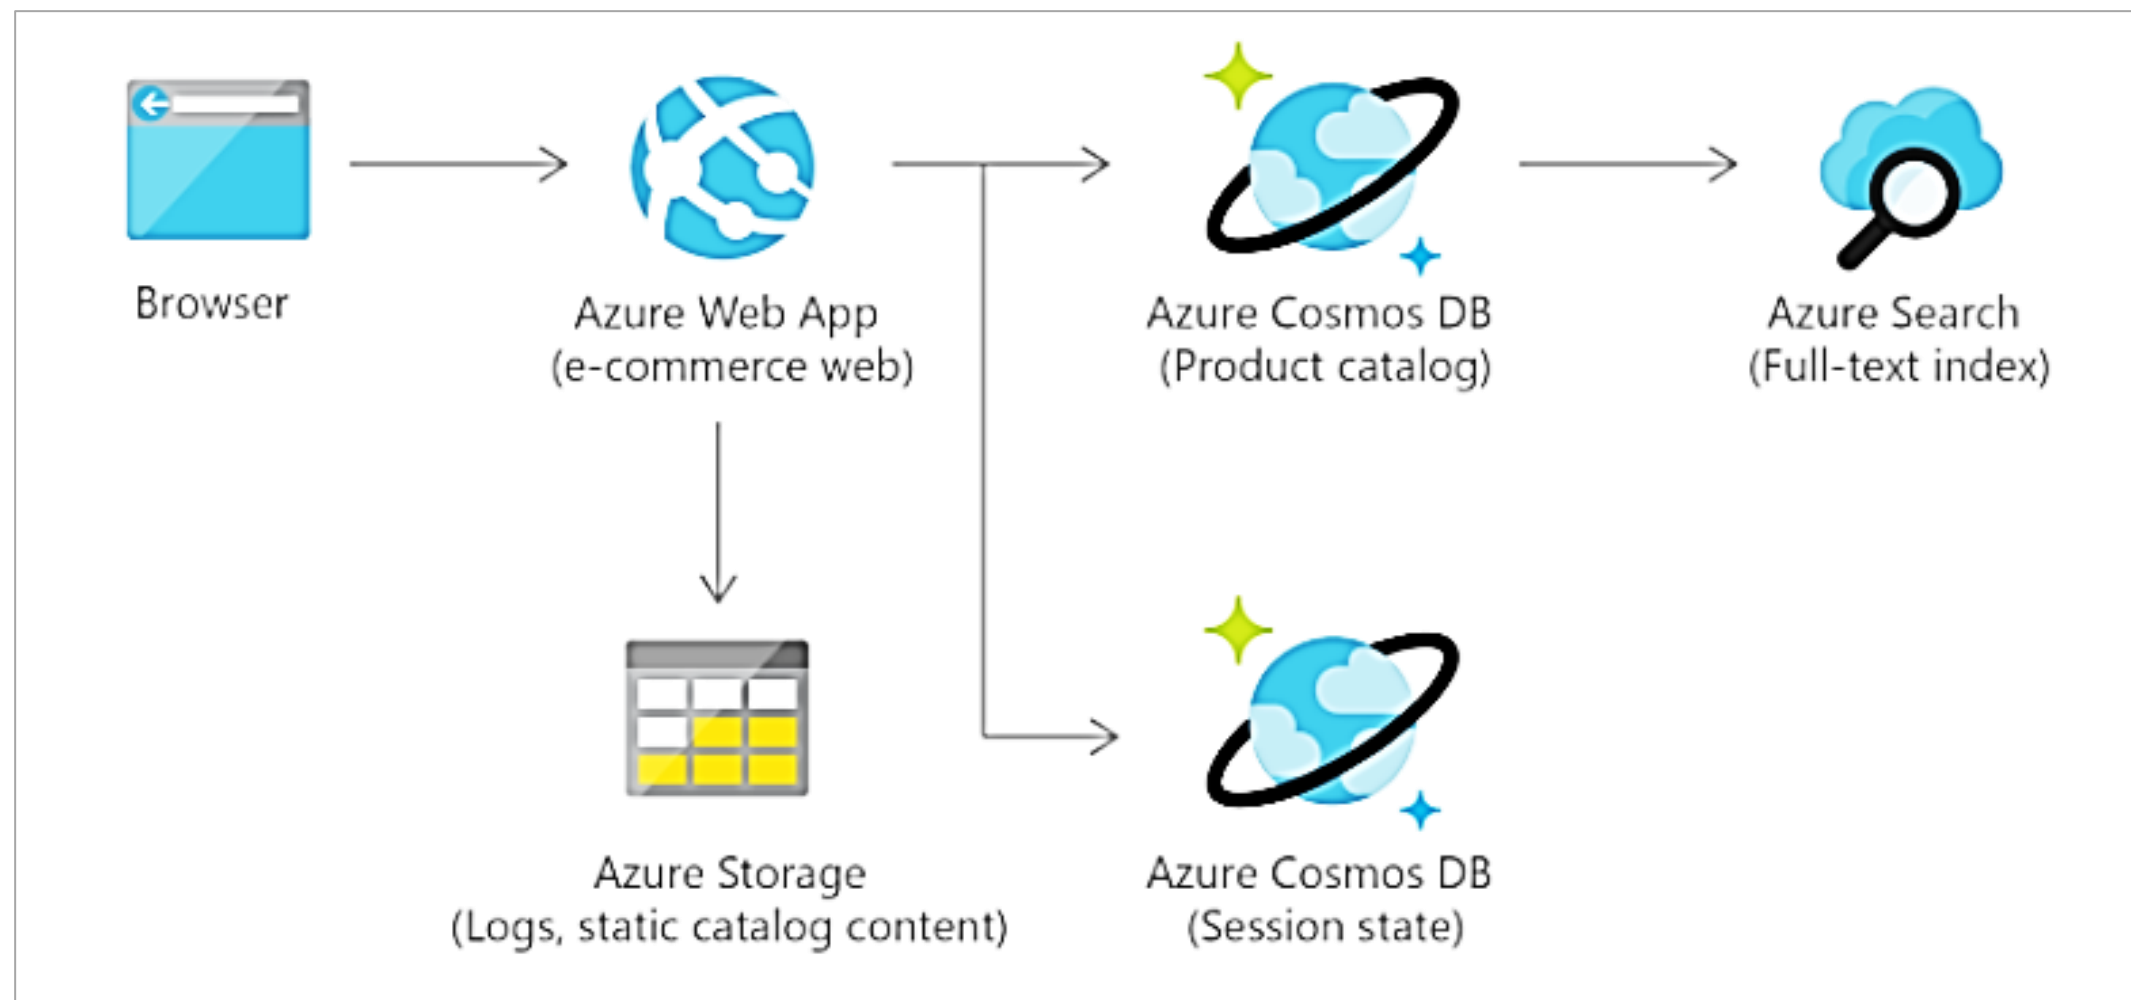
Modern games perform graphical processing on mobile or console clients but rely on the cloud for in-game content.





# Use Case of Azure Cosmos DB

Microsoft's e-commerce sites, such as the Windows Store and Xbox Live, rely heavily on Azure Cosmos DB.



# Azure Cosmos DB Table

The features of Azure table storage and Azure Cosmos Table API are:

Features	Azure Table storage	Azure Cosmos DB Table API
Latency	Fast but no upper bounds on latency	Single-digit millisecond latency for reads and writes
Query	Execution uses index for primary key and scans otherwise	Queries that can take advantage of automated indexing on properties
Indexing	Only primary index on Partition Key and Row Key	Automatic indexing on all properties by default without index management
Global distribution	Single region with one optional readable secondary read region for high availability	Turnkey global distribution from one to any number of regions

# Azure Cosmos DB Table

Features	Azure Table storage	Azure Cosmos DB Table API
Throughput	Scalability limit of 20,000 operations per second	Highly scalable with dedicated reserved throughput per table that is backed by SLAs and Accounts with no upper limit
Consistency	Strong within primary region and Eventual within secondary region	Five well-defined consistency levels to trade off availability, latency, throughput, and consistency based on needs.
SLAs	99.99% availability, depending on the replication strategy	Comprehensive SLAs covering availability, latency, throughput, and consistency.
Pricing	Consumption-based	Available in both consumption-based and provisioned capacity modes

# Storage Management Tools

# Manage Tiered Storage Using Azure Tools


There are various tools available to manage Azure Storage:




Azure Portal



Azure Storage Explorer



Azure CLI



Azure PowerShell

# Manage Tiered Storage Using Azure Tools

The following commands are used to manage the Azure tools:

```
az storage blob upload  
az storage blob list  
az storage blob download  
az storage blob set-tier
```

```
$Storage= "StorageAccountName"  
$Key = "Storage Account Key"  
$Container = "Blob Container"  
$Blobs = Get-AzureStorageblob -Container $Container  
$blob.icloudblob.setstandardblobtier("Cool")
```

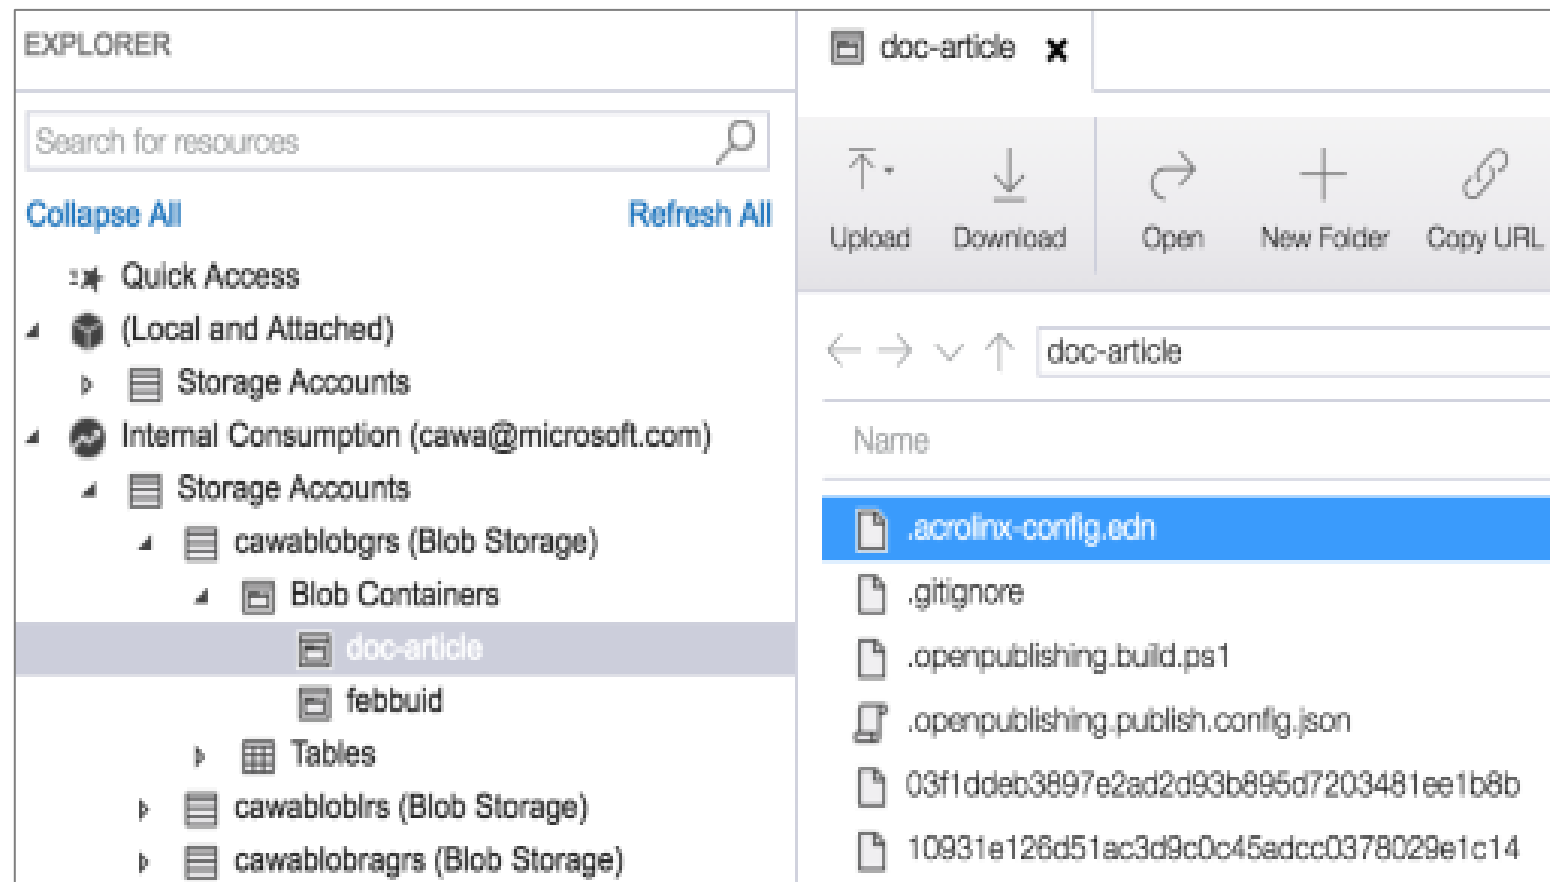
# Comparing Tools for Working with Azure Storage

The table given below shows the comparison of tools that can run on various platforms:

Azure Storage Client Tool	Supported Platform	Block Blob	Page Blob	Append Blob	Tables	Queues	Files
Azure Portal	Web	Yes	Yes	Yes	Yes	Yes	Yes
Azure Storage Explorer	Windows, OSX	Yes	Yes	Yes	Yes	Yes	Yes
Microsoft Visual Studio Cloud Explorer	Windows	Yes	Yes	Yes	Yes	Yes	No

# Azure Storage Explorer

Azure Storage Explorer is a single software for Windows, macOS, and Linux that makes working with Azure Storage data simple.

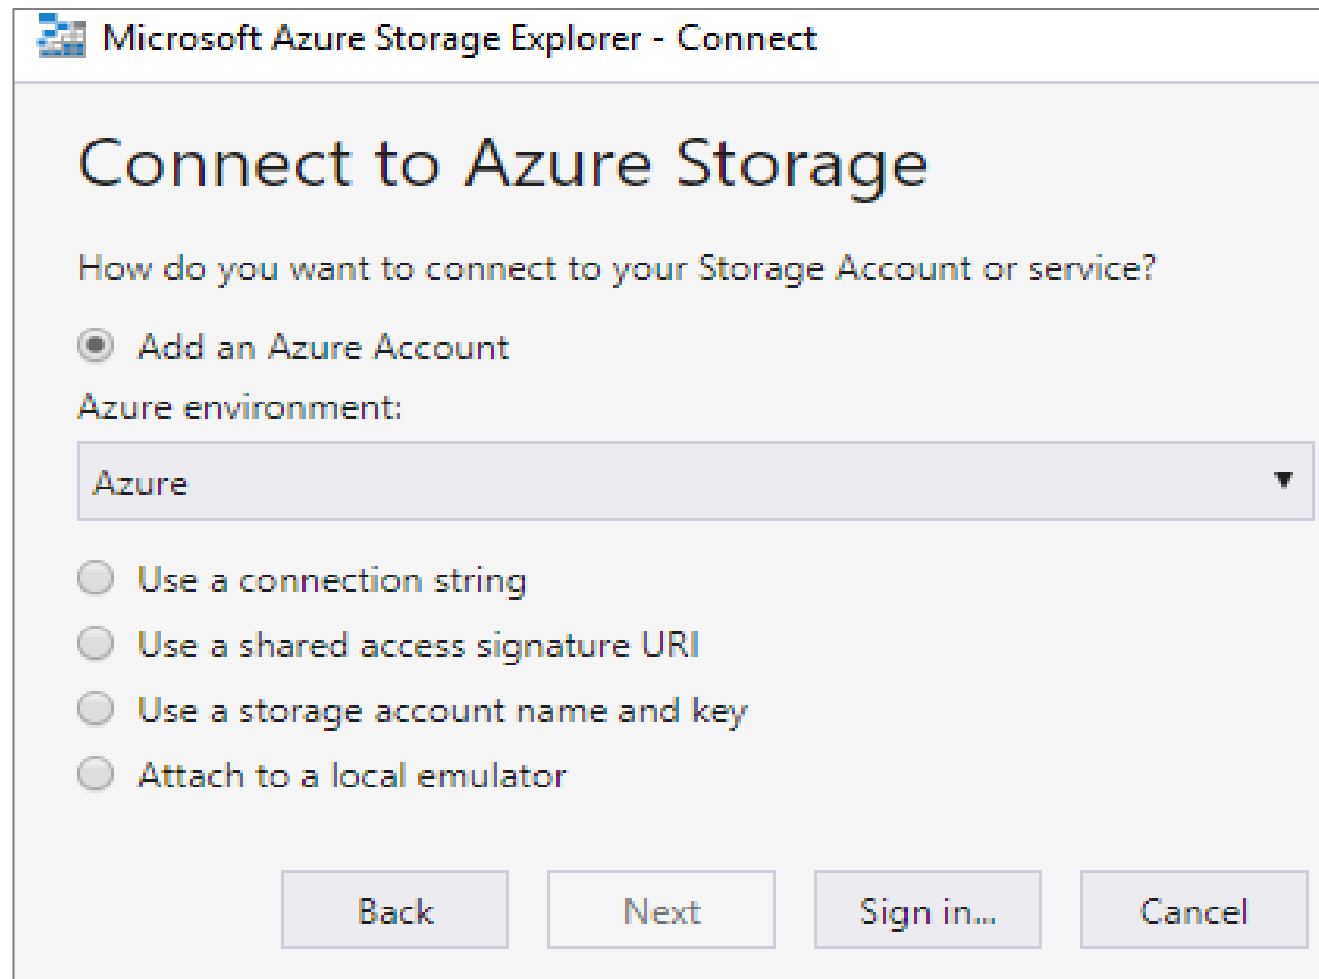


- Access multiple accounts and subscriptions
- Create, delete, view, and edit storage resources
- View and edit Blob, Queue, Table, File, Cosmos DB storage, and Data Lake Storage
- Obtain shared access signature (SAS) keys
- Available for Windows, Mac, and Linux



# Azure Storage Connection

The procedure for connecting to Azure storage is given below:



The screenshot shows the 'Microsoft Azure Storage Explorer - Connect' window. The title bar reads 'Microsoft Azure Storage Explorer - Connect'. The main heading is 'Connect to Azure Storage'. Below this, a question asks 'How do you want to connect to your Storage Account or service?'. There are four radio button options: 'Add an Azure Account' (which is selected), 'Use a connection string', 'Use a shared access signature URI', and 'Attach to a local emulator'. Under the 'Add an Azure Account' option, there is a label 'Azure environment:' followed by a dropdown menu currently showing 'Azure'. At the bottom of the window, there are four buttons: 'Back', 'Next', 'Sign in...', and 'Cancel'.

## Options

- Connect to an Azure subscription
- Work with local development storage (emulator)
- Attach to external storage
- Attach a storage account or storage service by using a shared access signature
- Connect to an Azure Cosmos DB account by using a connection string

# Assisted Practice

## Managing Cosmos DB

Duration: 10 Min.

### Problem Statement:

You need to create a Cosmos DB account and identify different options for deploying and configuring Cosmos DB. You also need to add some data to the Cosmos DB to check its functionality further by running queries.

# Assisted Practice: Guidelines

Steps to manage storage access keys:

1. Creating a Cosmos DB account
2. Creating a database
3. Adding some data to the cosmos DB database
4. Querying the data using SQL APIs



## Design solution for Database

Duration: 25 min.




**Project Agenda:** Design a solution for database

**Description:** You have been given a project to move a database from on-premise to the cloud. The data is in a document format. You need to choose the right database option. You also need to ensure that you have support for querying items using the Structured Query Language (SQL) syntax. Suggest a solution for this migration.

**Perform the following:**

1. Creating an Azure Cosmos DB account by selecting Core (SQL) API option.



**Thank you**